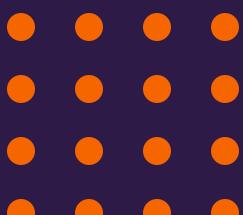




# Final Exam Review

Aaron J. Masino, PhD

Associate Professor





## Exam logistics

- Location: Academic Success Center, Room 118
- Date / time: May 2<sup>nd</sup>, 7:00-9:30 PM
- Mode: Canvas (**Bring your laptop!!!**)
- The test is closed-book. Internet and AI resources are not allowed.
- Python terminal (IPython), calculator, calculator app may be used.

# Coding?

- There are NO coding problems (i.e., you will NOT need to write code)
- There ARE code interpretation problems
- Review the labs

Example: What is the value of z in the code below?

```
x = numpy.array([1, 2, 3])
y = numpy.array([5, 1, 0])
z = x + y
```

- A. [1, 2, 3, 5, 1, 0]
- B. [[1, 2, 3], [5, 1, 0]]
- C. [6, 3, 3]



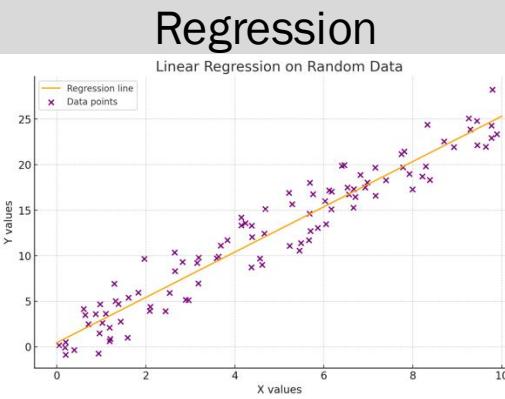
# Python programming

- Data structures
  - Core Python: lists, sets, dictionaries
  - Numpy Arrays
  - PyTorch Tensors
  - Pandas DataFrame
- Index slicing for lists, numpy arrays, PyTorch Tensor
  - How many elements in  $l[x:x+7]$
- Numpy and PyTorch broadcasting
- PyTorch matrix operations: transpose, multiplication

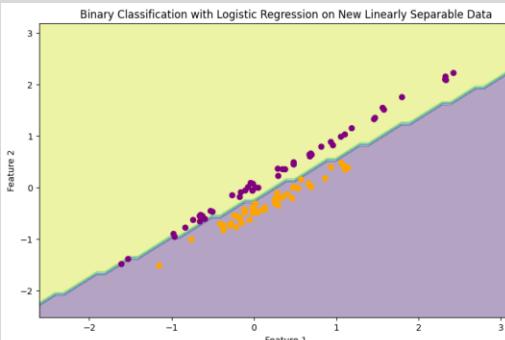
# Machine learning concepts

- Supervised vs. unsupervised learning.
  - Which one requires labeled outcomes for training?
  - Tasks

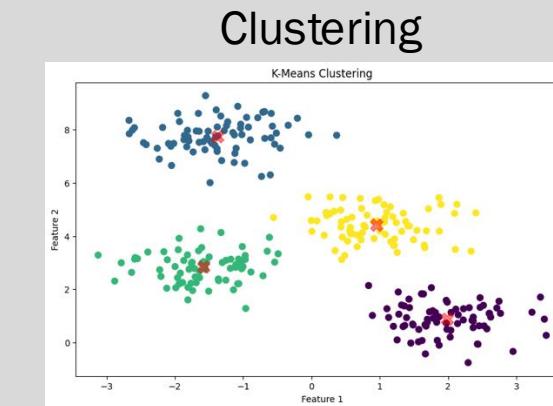
Supervised



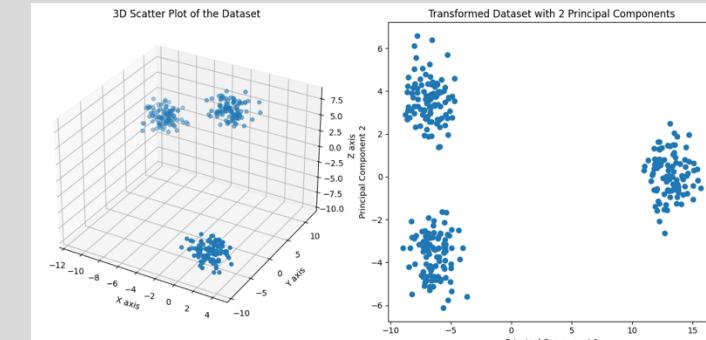
Classification



Unsupervised

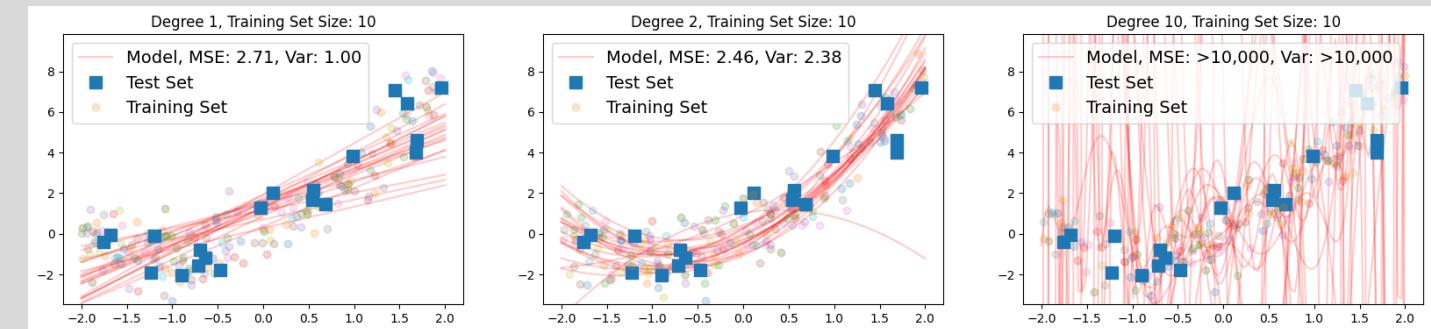
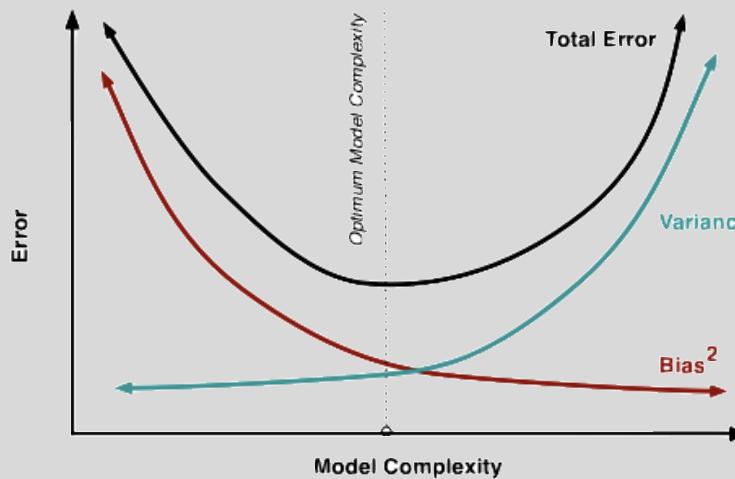


Dimensionality Reduction



# Machine learning concepts

- Bias vs variance



- Model selection
  - K-fold CV
  - Bootstrapping
  - Train / validation / test



# Linear Regression

- Model form

$$\hat{y} = \beta_0 + \beta_1 * age + \beta_2 * height + \beta_3 * weight$$

- Assessing model fit

- **R<sup>2</sup>** – the proportion of explained variance. Range [0, 1]. Higher is better.  
What happens if we add predictor variables?
- **F-statistic** – used to assess likelihood of the null hypothesis

$$H_0 : \beta_1 = \beta_2 = \cdots = \beta_p = 0$$

- **Coefficient standard error** – average deviation of coefficient estimate from the actual coefficient
- **t-statistic** - For a given coefficient, it is a test of the null hypothesis that the individual coefficient is zero

# Linear Regression

## Python StatsModels library

Summary data for multiple linear regression:								
OLS Regression Results								
Dep. Variable:	PickupCount			R-squared:	0.382			
Model:	OLS		Adj. R-squared:	0.380				
Method:	Least Squares			F-statistic:	205.0			
Date:	Fri, 27 Sep 2024		Prob (F-statistic):	1.58e-103				
Time:	14:57:21		Log-Likelihood:	-4131.3				
No. Observations:	1000			AIC:	8271.			
Df Residuals:	996			BIC:	8290.			
Df Model:	3							
Covariance Type:	nonrobust							
	coef	std err	t	P> t	[0.025	0.975]		
const	41.3965	1.990	20.801	0.000	37.491	45.302		
x1	-9.2296	0.714	-12.928	0.000	-10.631	-7.829		
x2	1.0189	0.069	14.735	0.000	0.883	1.155		
x3	-0.0267	0.002	-14.090	0.000	-0.030	-0.023		

What do these numbers mean?

Our model:

$$\hat{y} \approx 41.40 - 9.23 * x1 + 1.02 * x2 - 0.03 * x3$$

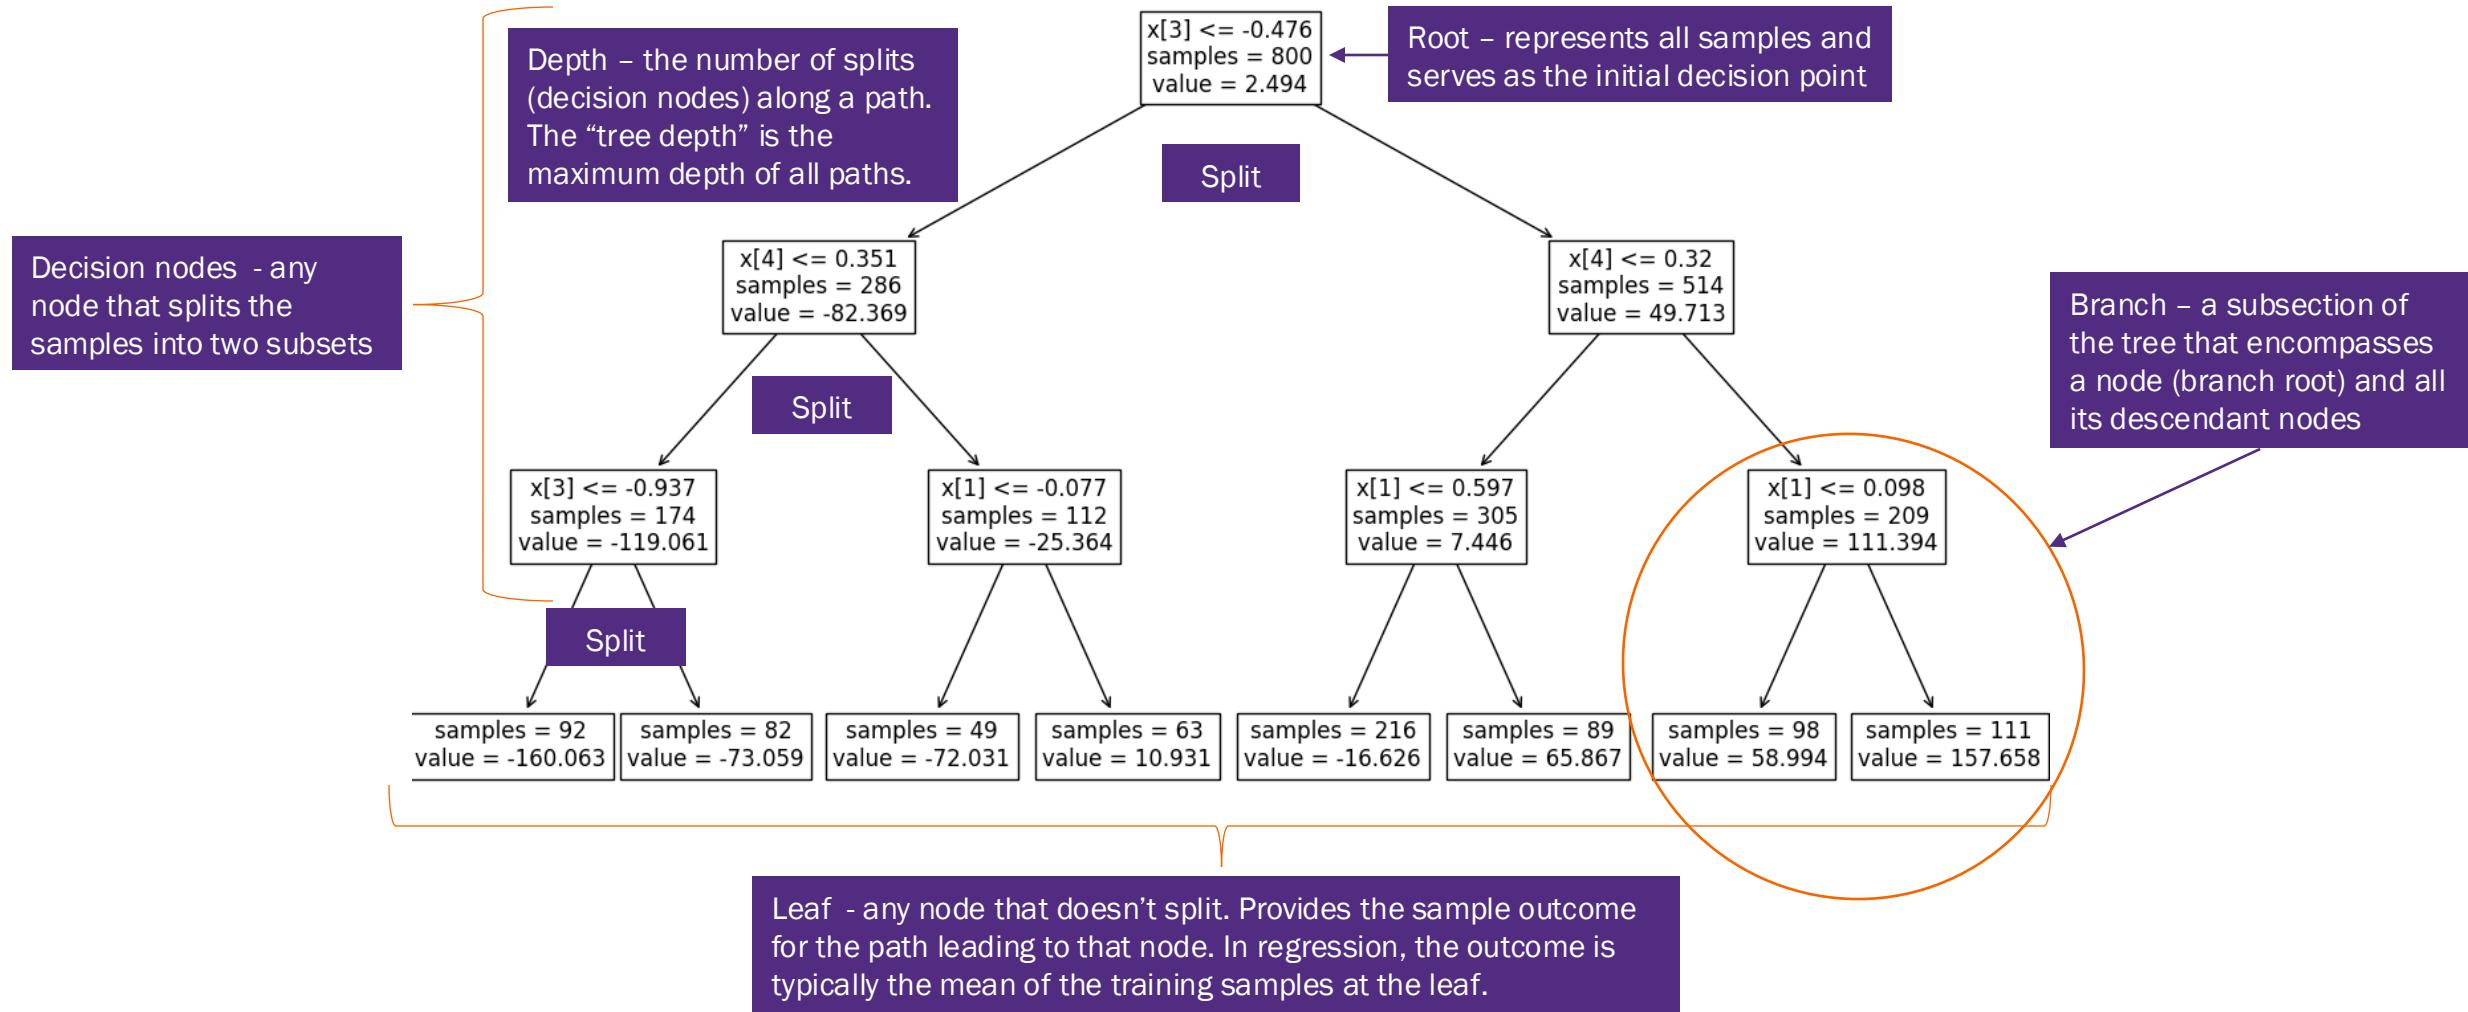
# Logistic Regression

- Binary classification model

$$P(Y = 1) = \frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}} = \frac{1}{1 + e^{-(\beta_0 + \beta_1 X)}}$$

- We classify all observations for which  $\hat{P}(Y = 1) \geq \gamma$  to be in the group associated with  $Y = 1$  and all observations for which  $\hat{P}(Y = 0) < \gamma$  to be in the group associated with  $Y = 0$ . We select  $0 < \gamma < 1$ .

# Decision tree terminology



# Tree Ensembles

## Bagging Decision Trees

- How can bagging be used with decision trees for regression?
1. Construct  $B$  bootstrapped datasets from available training data
  2. For each bootstrapped dataset, construct a deep, unpruned tree
  3. For a given test sample, the final output is the average of the  $B$  trees for that sample
- ISSUE: trees are highly correlated

## Random Forests

- Extension of bagging that decorrelates the trees
- Follows the same procedure as bagging, except that for each split, only  $m < p$  predictors are considered
- As with normal bagging, the number of trees (bootstrapped samples) does not affect variance

## Boosting Trees

- Similar to bagging in that multiple trees are created
- No bootstrapping
- Instead, we fit a sequence of small trees to the residuals  $r_i = y_i - \hat{y}_i$  (initially all  $r_i = y_i$ )
- Unlike bagging, boosting variance is sensitive to the value of  $B$ . Creating too many trees can lead to overfitting.

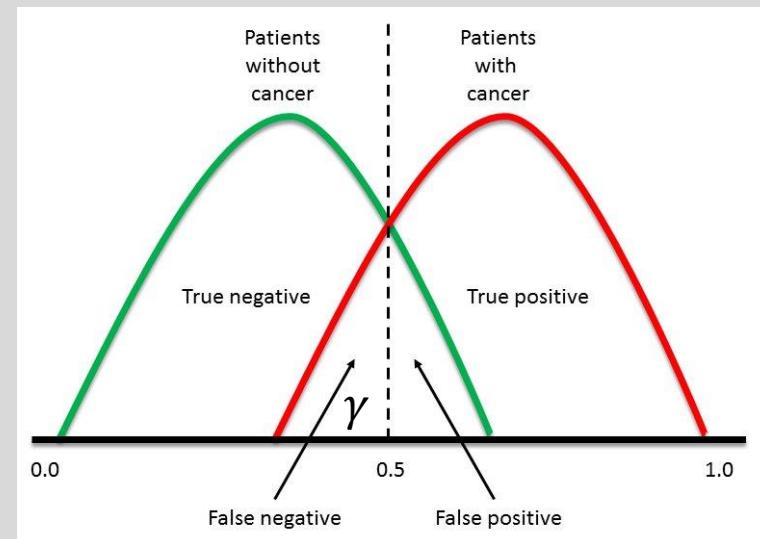
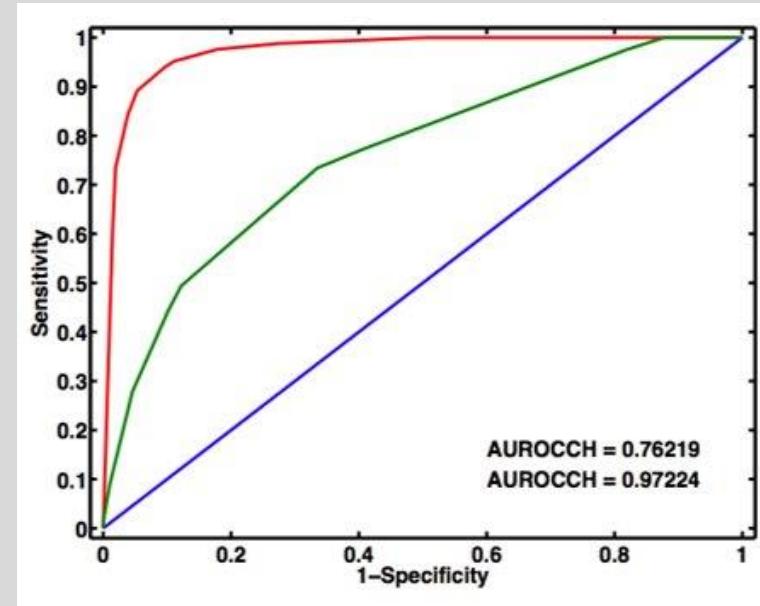
# Classification Performance Metrics

		Actual Class	
		Positive	Negative
Predicted Class	Positive	True Positive (TP)	False Positive (FP) <i>False Alarm Type I Error</i>
	Negative	False Negative (FN) <i>Type II Error</i>	True Negative (TN)

- Accuracy =  $(TP+TN)/(TP+TN+FP+FN)$
- Sensitivity (a.k.a Recall, True Positive Rate) =  $TP/(TP+FN) = TP/P$ 
  - The fraction of actually positive samples predicted positive by the model
- Specificity (a.k.a. True Negative Rate) =  $TN/(TN+FP) = TN/N$ 
  - The fraction of actually negative samples predicted negative by the model
- Precision (a.k.a. Positive Predictive Value) =  $TP/(TP+FP)$ 
  - The fraction of samples predicted positive by the model that are actually positive
- False Positive Rate (a.k.a. Probability of false alarm) =  $FP/(FP+TN) = FP/N$
- F1 Score =  $2(\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$
- Balanced Accuracy =  $(\text{Sensitivity} + \text{Specificity})/2$

# Classifier performance with ROC and AUC

- Recall, we can select  $0 < \gamma < 1$  as a threshold s.t.  $\hat{P}(Y = 1) \geq \gamma$  implies  $\hat{y} = 1$ . What happens if we change  $\gamma$ ?
- Receiver Operating Characteristic Curve (ROC)
  - Illustrates tradeoff between FPR and TPR of a binary classifier as the discrimination threshold,  $\gamma$ , is adjusted
  - Axes
    - $1 - \text{Specificity}$  (false positive rate) =  $1 - \text{TN}/\text{N}$
    - Sensitivity (true positive rate) =  $\text{TP}/\text{P}$
- Area under the curve (AUC)
  - Summary statistic
  - AUC = 0.5 indicates no predictive value (random guessing)
  - AUC = 1.0 indicates perfect predictive value





# K-means clustering

- Given  $n$  samples partition the samples into  $K$  distinct groups,  $C_1, \dots, C_k$
- Each sample is assigned to exactly one group (cluster)
- Conceptual approach is to assign cluster membership that minimizes

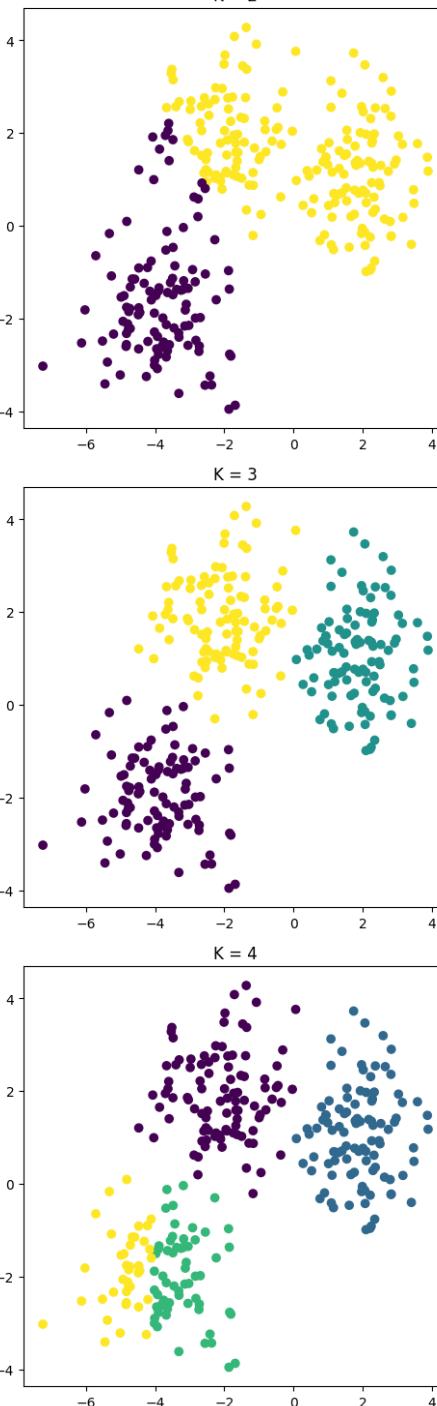
$$\sum_{k=1}^K W(C_k)$$

where  $W(C_k)$  is a measure of *intra-cluster variation*

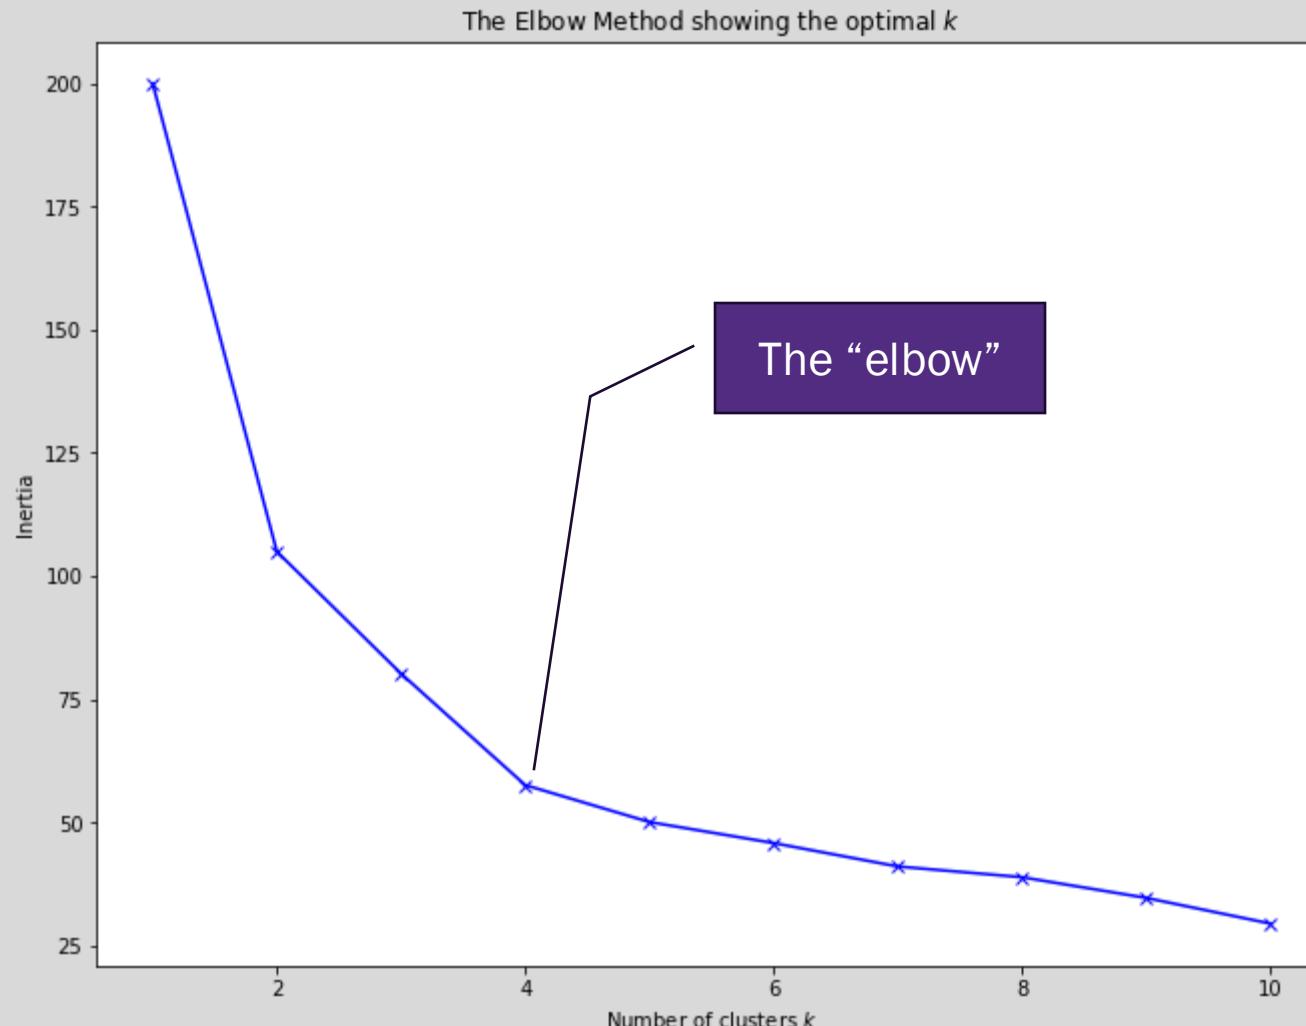
- Letting  $W(C_k)$  be the squared Euclidean distance, we seek  $C_1, \dots, C_k$  that minimize

$$\sum_{k=1}^K \frac{1}{|C_k|} \sum_{i,i' \in C_k} \sum_{j=1}^p (x_{i,j} - x_{i',j})^2$$

- Other metrics: Hamming (binary vectors), Manhattan (integer vectors), Gower's (combined binary, numerical, categorical)



# Elbow Method – How many clusters?



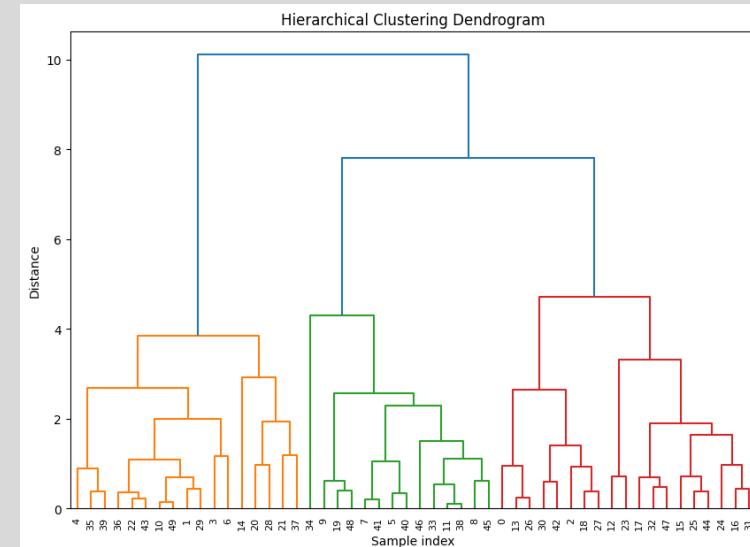
# Agglomerative Clustering

The algorithm:

- Each observation starts as its own cluster
- At each step of the algorithm, **two clusters** that are most “similar” are combined into a new larger cluster
- This process of combining clusters is repeated until all observations are members of one single large cluster

## Linkage functions

- Complete (or maximum) Linkage Clustering: For two clusters, determine the maximum dissimilarity between any observation in the first cluster and any observation in the second cluster.
- Single Linkage Clustering: For two clusters, determine the minimum dissimilarity between any observation in the first cluster and any observation in the second cluster.
- Average Linkage Clustering: Compute all pairwise dissimilarities between observations in the first and second cluster and calculate the average.





# Feed forward neural networks

- This layer is fully connected
- Hidden layer of multiple computation nodes

$$z_j^{(1)} = \left[ \sum_{j=0}^n w_{j,k}^{(1)} x_j \right]$$

$$a_j^{(1)} = h(z_j^{(1)})$$

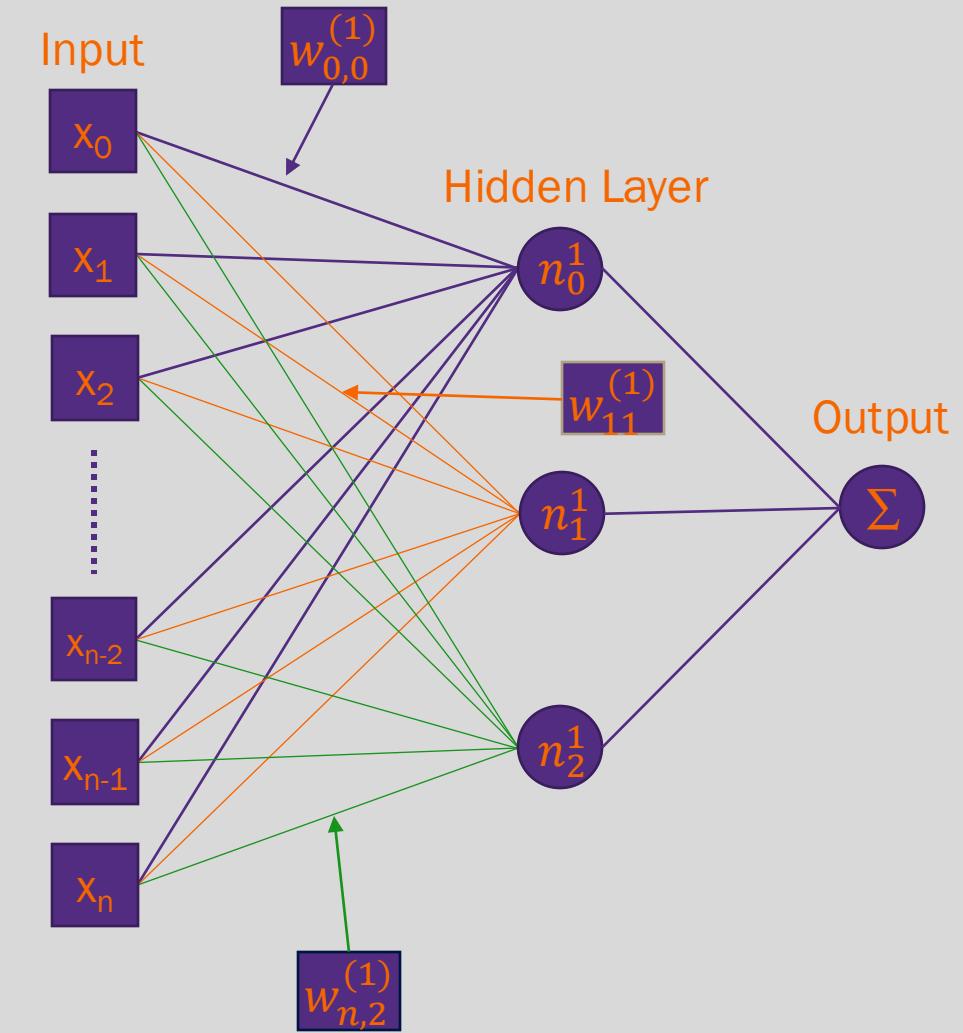
- Output layer is a weighted linear sum

$$z^{(out)} = \left[ \sum_{j=0}^k w_j^{(2)} a_j^{(1)} \right]$$

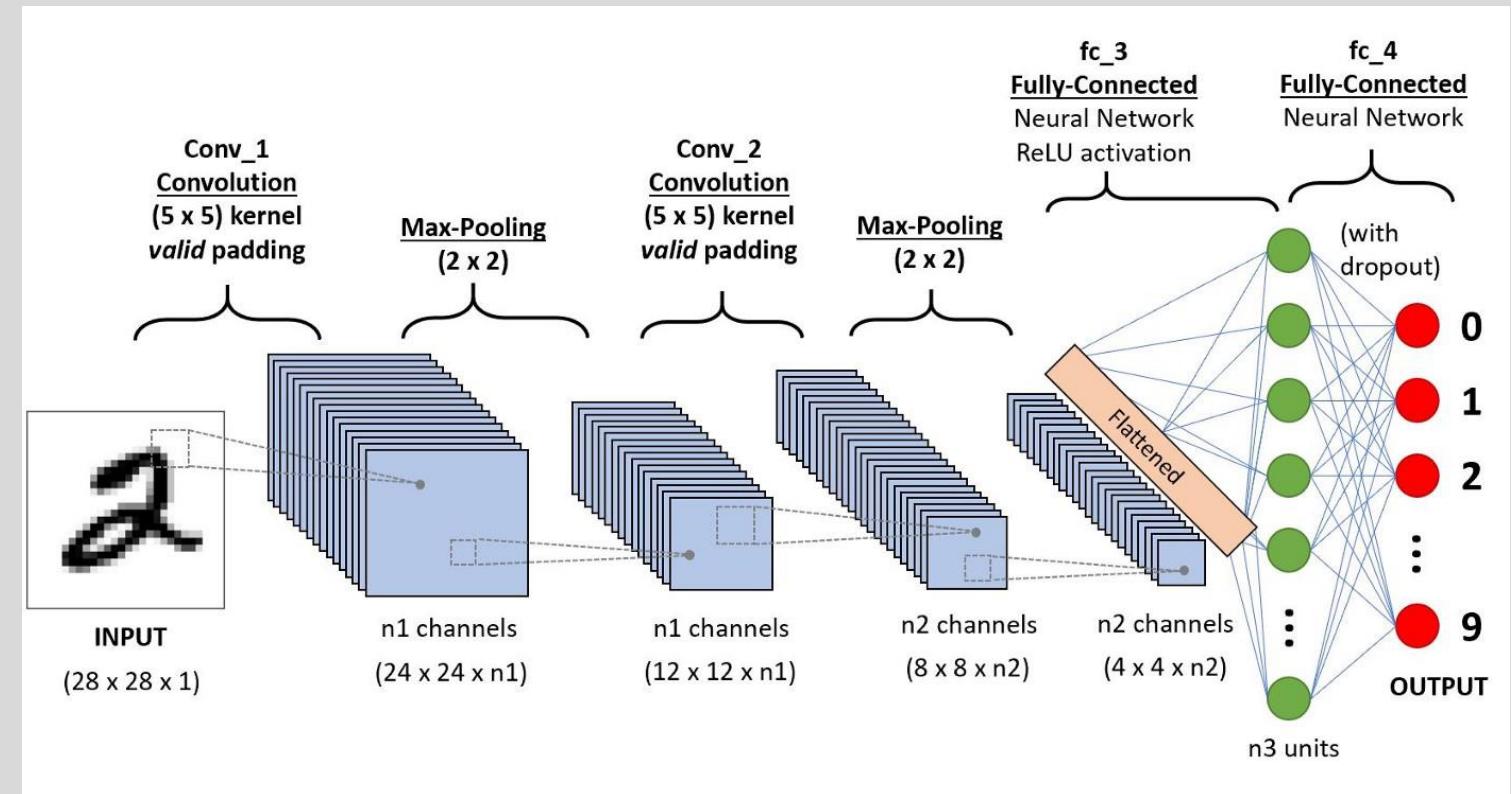
- In classification, sigmoid (binary) or softmax (multiclass) function is applied to output

$$\sigma(z) = \frac{1}{1+e^{-z}} \quad \sigma_{SM}(z_k, z) = \frac{e^{z_k}}{\sum_{j=1}^M e^{z_j}}$$

How many parameters are in this model?

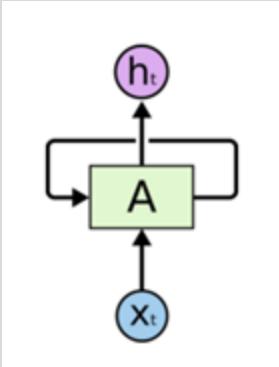


# Convolutional Neural Networks



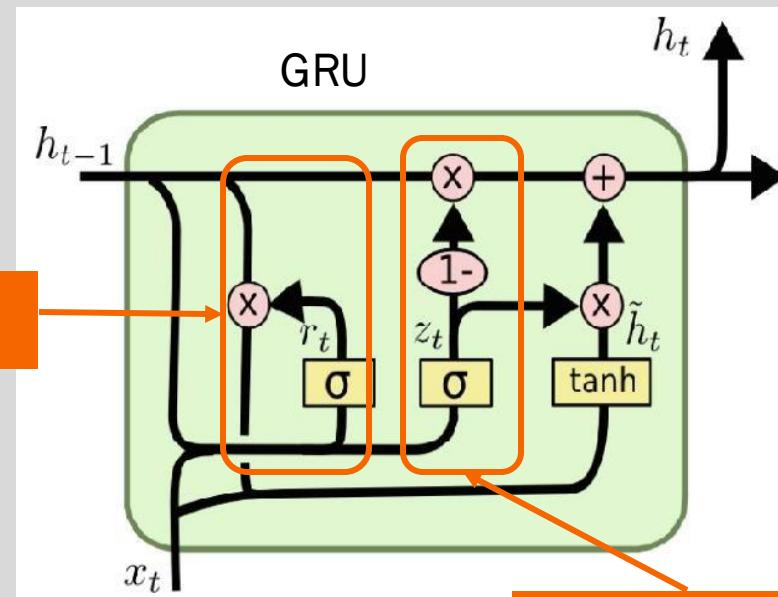
How many inputs are there from the last convolution layer to the fully connected layer?

# Recurrent Neural Networks?



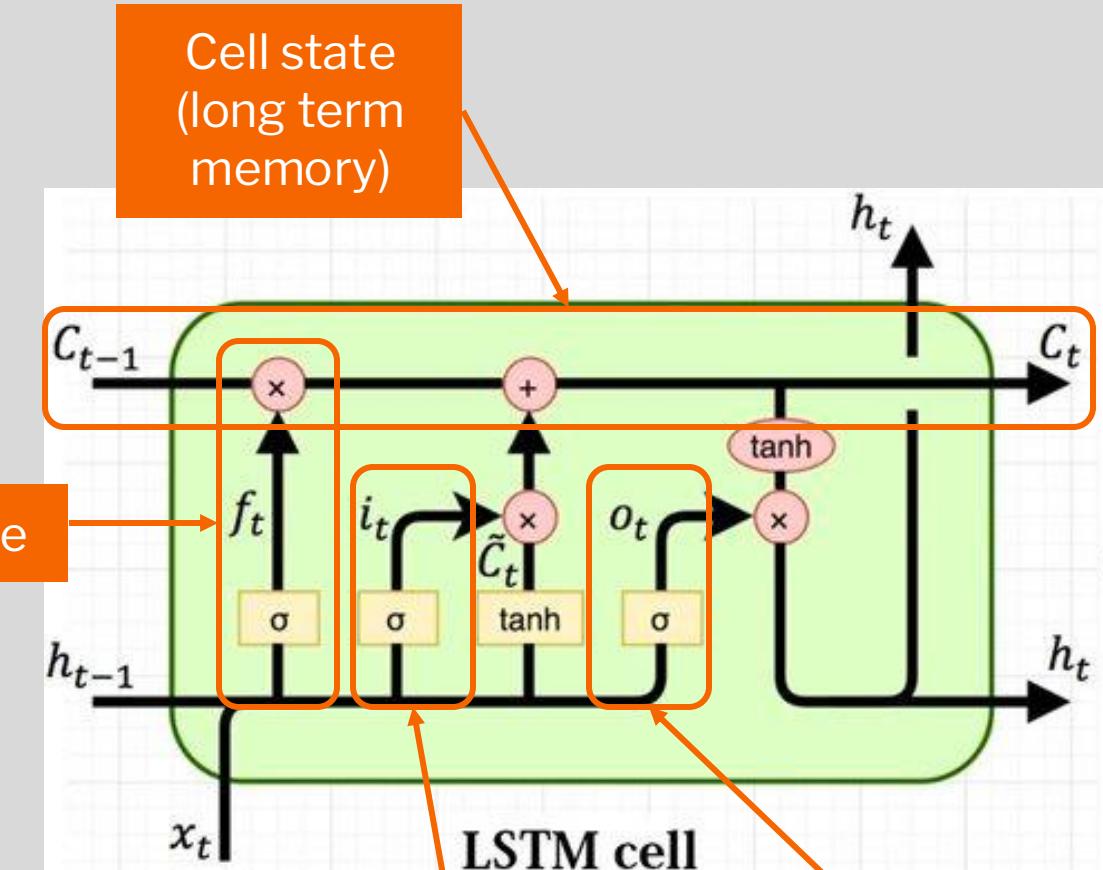
RNN's have self-loops  
What do these gates do?

Forget gate



Reset Gate

Update gate



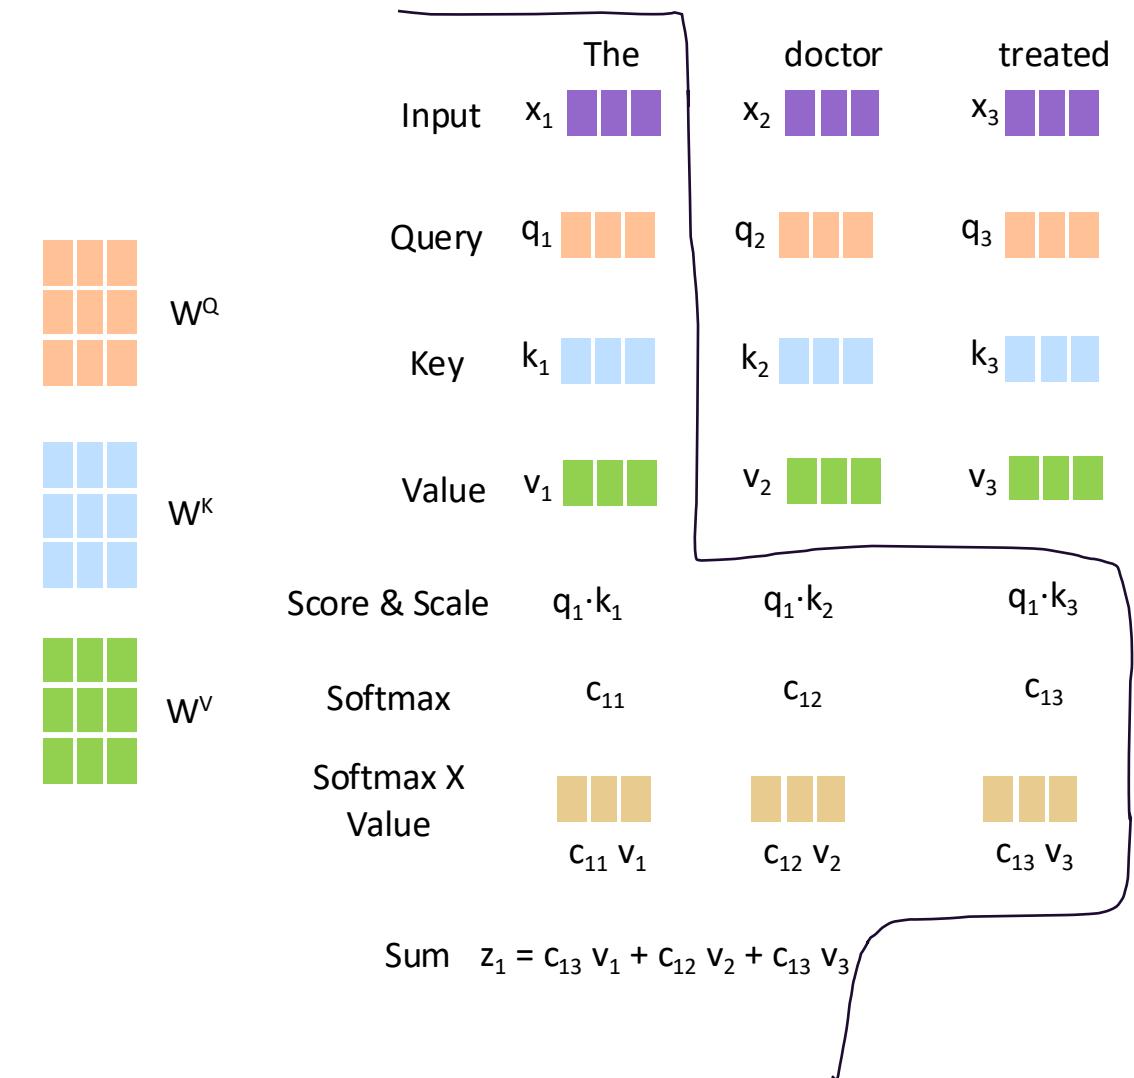
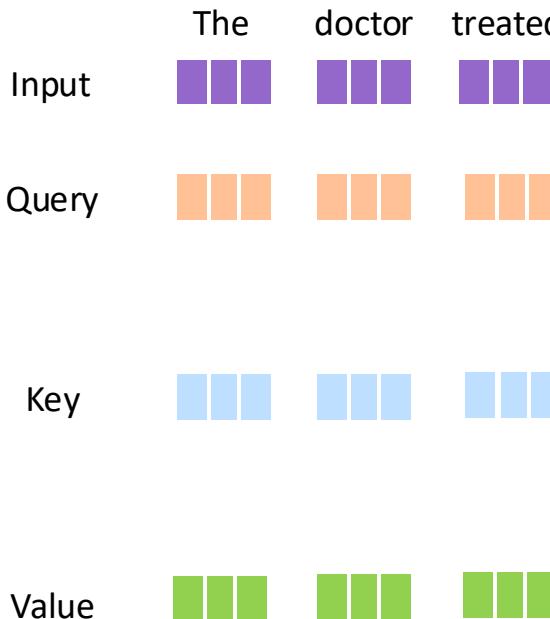
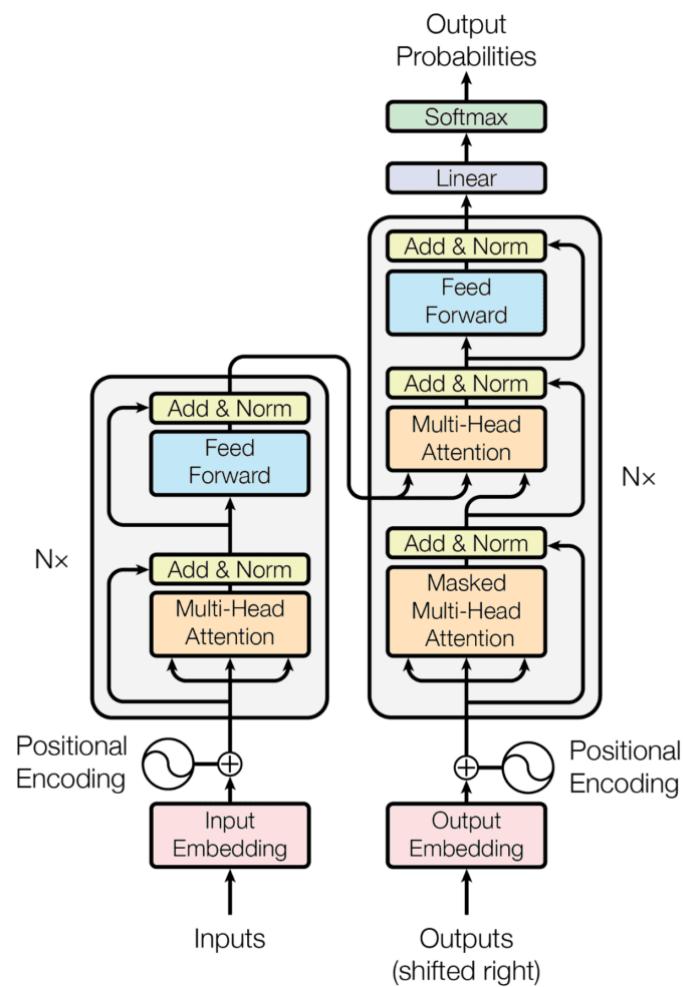
Cell state  
(long term  
memory)

Input gate

Output gate

LSTM cell

# Transformers and Attention



# Deep-learning implementation (PyTorch)

## Dimensions

### Dimensions

#### Dimensions

```
class CIFAR10Output(nn.Module):
    def __init__(self, input_dim, num_classes):
        super().__init__()
        self.flatten = nn.Flatten() # Flatten the input tensor
        self.linear = nn.Linear(input_dim, num_classes) # Linear layer with 10 output units

    def forward(self, x):
        x = self.flatten(x) # Flatten the input tensor
        x = self.linear(x) # Pass through the linear layer
        return x

# Test the module with random input data
decoder = CIFAR10Output(num_kernels*15*15, 10)
input_tensor = torch.randn(5, num_kernels, 15, 15) # Batch size of 5, input tensor shape [5, 32, 7, 7]
rslt = decoder(input_tensor)
print("Rslt shape:", rslt.shape) # Expected output shape: [5, 10]
print(rslt.dtype)
```

