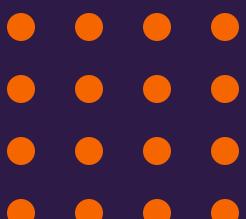




# Introduction to Convolutional and Recurrent Neural Networks

**Aaron J. Masino, PhD**  
Associate Professor, School of Computing



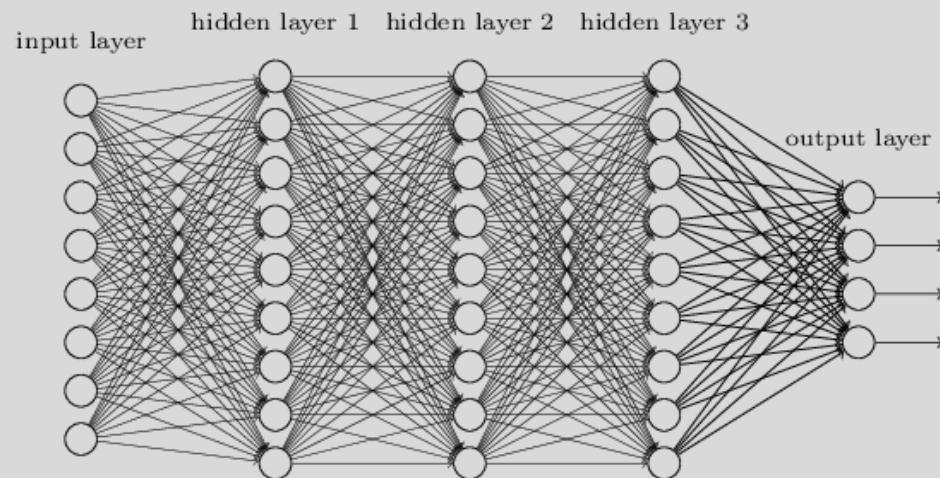
# Outline

- Convolutional Neural Network
- Recurrent Neural Network (RNN)
  - Vanilla RNN
  - LSTM
  - GRU



# Convolutional Neural Network (CNN)





## Motivations for CNNs

- Spatial invariant feature learning – some data types (e.g., images) have patterns that are informative regardless of their location
- Learning parameter efficiency
  - Variance (overfitting) – we generally prefer models with fewer parameters to mitigate overfitting
  - Training efficiency – models with fewer parameters tend to be easier to train
- Fully connected ANNs lack these properties

# Consider images

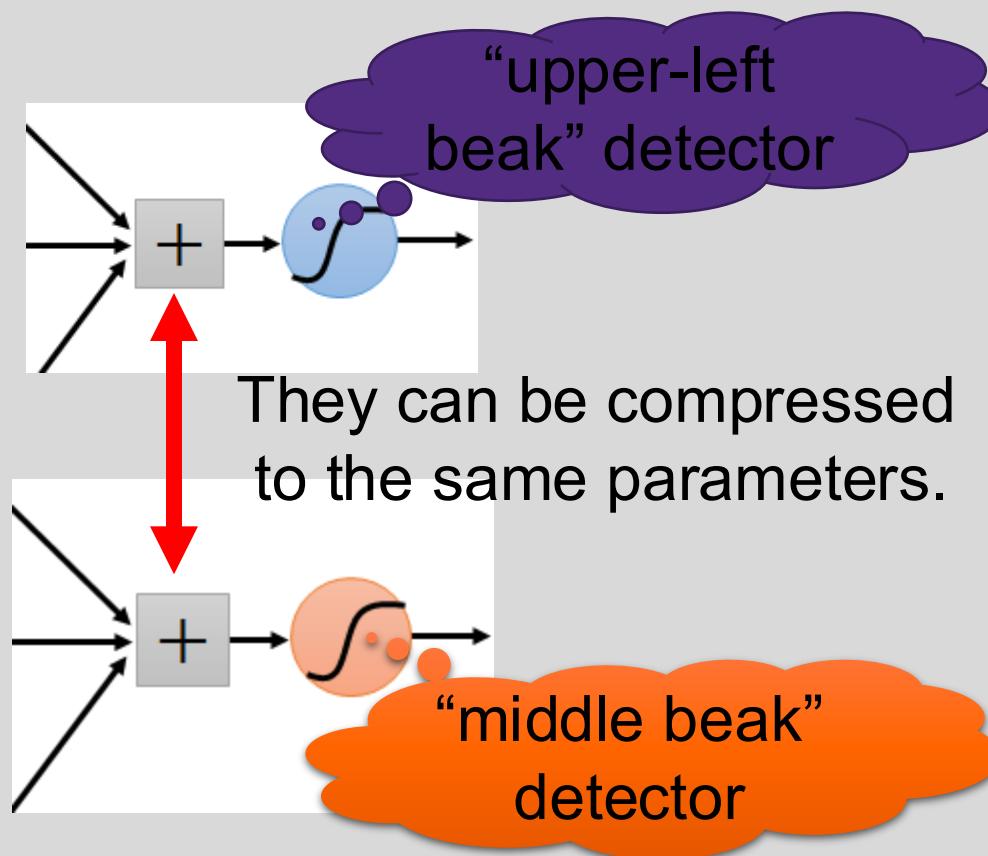
Some patterns are much smaller than the whole image

Can represent a small region with fewer parameters?



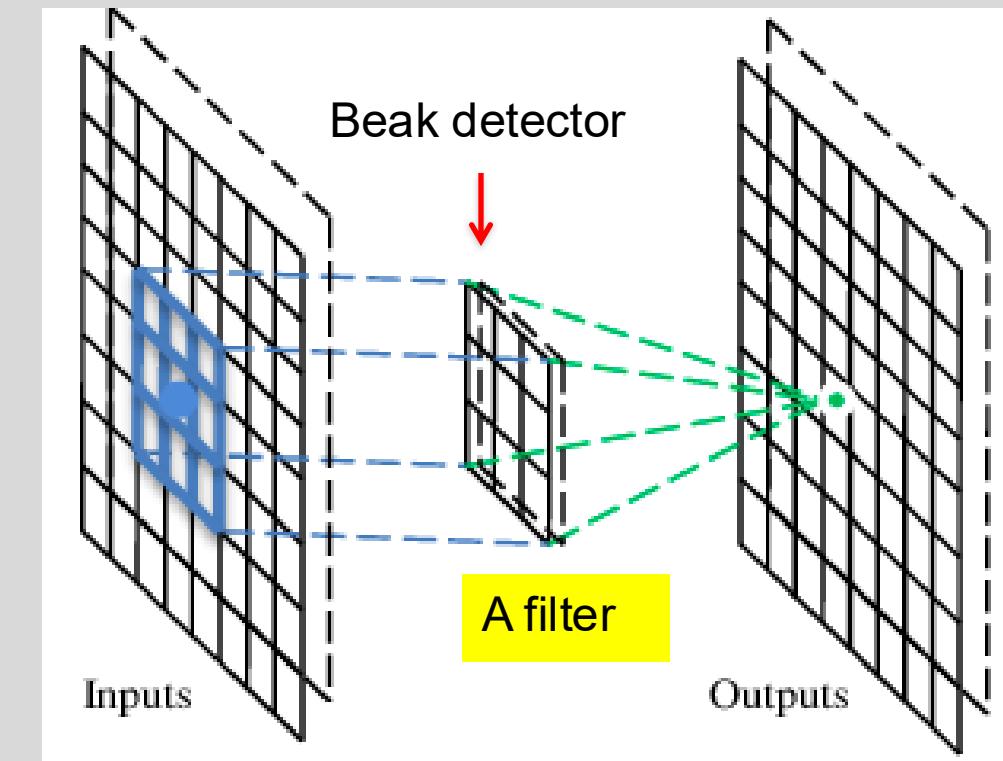
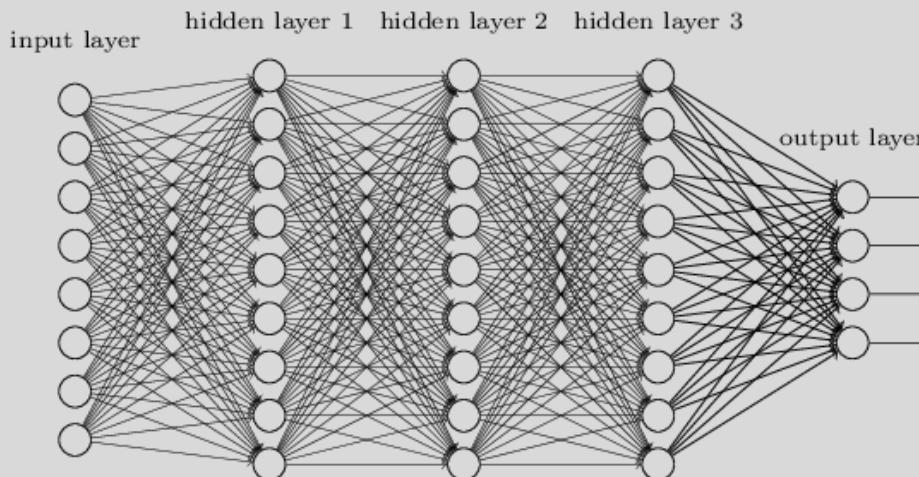
Same pattern appears in different places:  
They can be compressed!

Do we need a separate detector for every location?

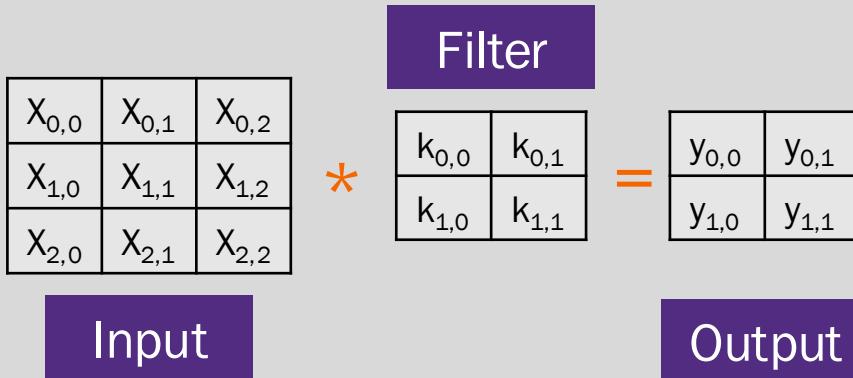


# Can we reuse our “detector”

- We want an architecture that will allow us to apply the detector over the entire input
- Fully connected layer does NOT do this
- Enter *convolution*



# Convolution Operator



Convolution: Filter,  $k$ , is convolved with input by sliding it across the input with stride,  $s$ . At each location of overlap, the dot product of the filter and input region is computed to from the output.

Output for 2 X 2 kernel with stride of 1

$$y_{0,0} = k_{0,0}x_{0,0} + k_{0,1}x_{0,1} + k_{1,0}x_{1,0} + k_{1,1}x_{1,1}$$

$$y_{0,1} = k_{0,0}x_{0,1} + k_{0,1}x_{0,2} + k_{1,0}x_{1,1} + k_{1,1}x_{1,2}$$

$$y_{1,0} = k_{0,0}x_{1,0} + k_{0,1}x_{1,1} + k_{1,0}x_{2,0} + k_{1,1}x_{2,1}$$

$$y_{0,0} = k_{0,0}x_{1,1} + k_{0,1}x_{1,2} + k_{1,0}x_{2,1} + k_{1,1}x_{2,2}$$

# Convolution

These are the network parameters to be learned.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

: :

Each filter detects a small pattern (3 x 3).

# Convolution

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

Dot  
product



3

-1

6 x 6 image

# Convolution

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

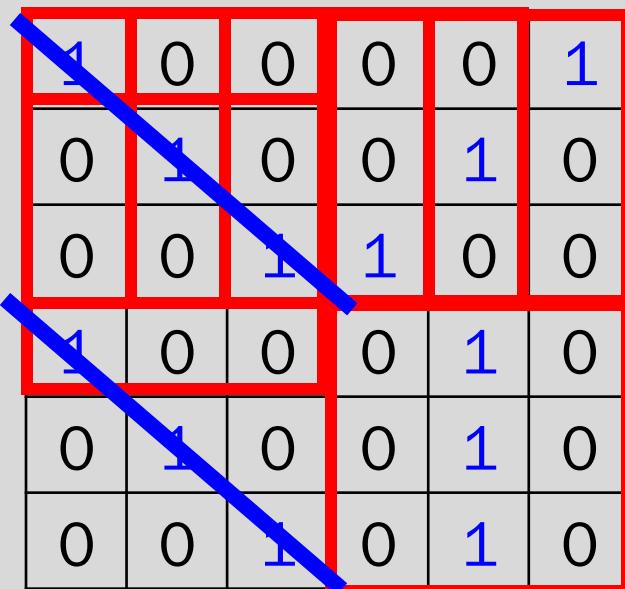
3

-3

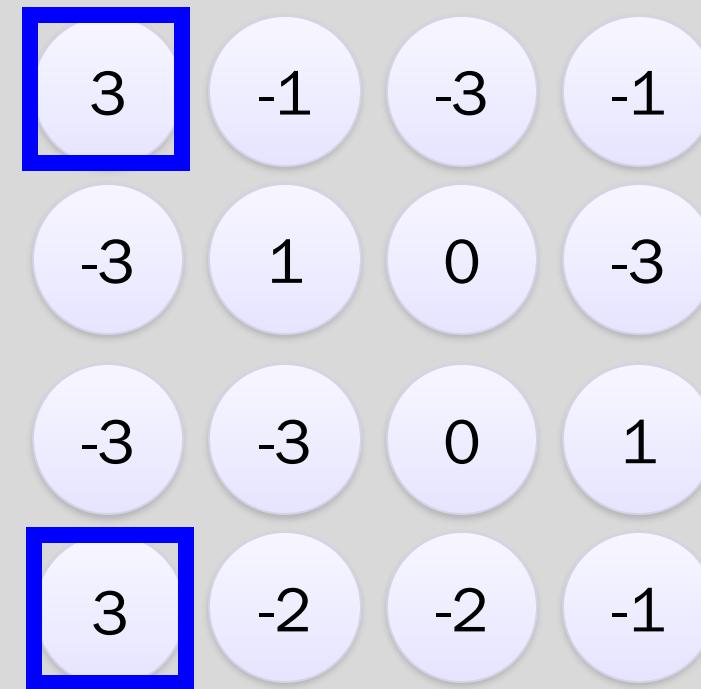
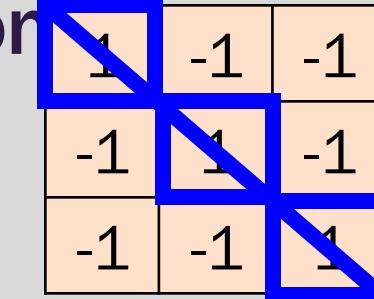
6 x 6 image

# Convolution

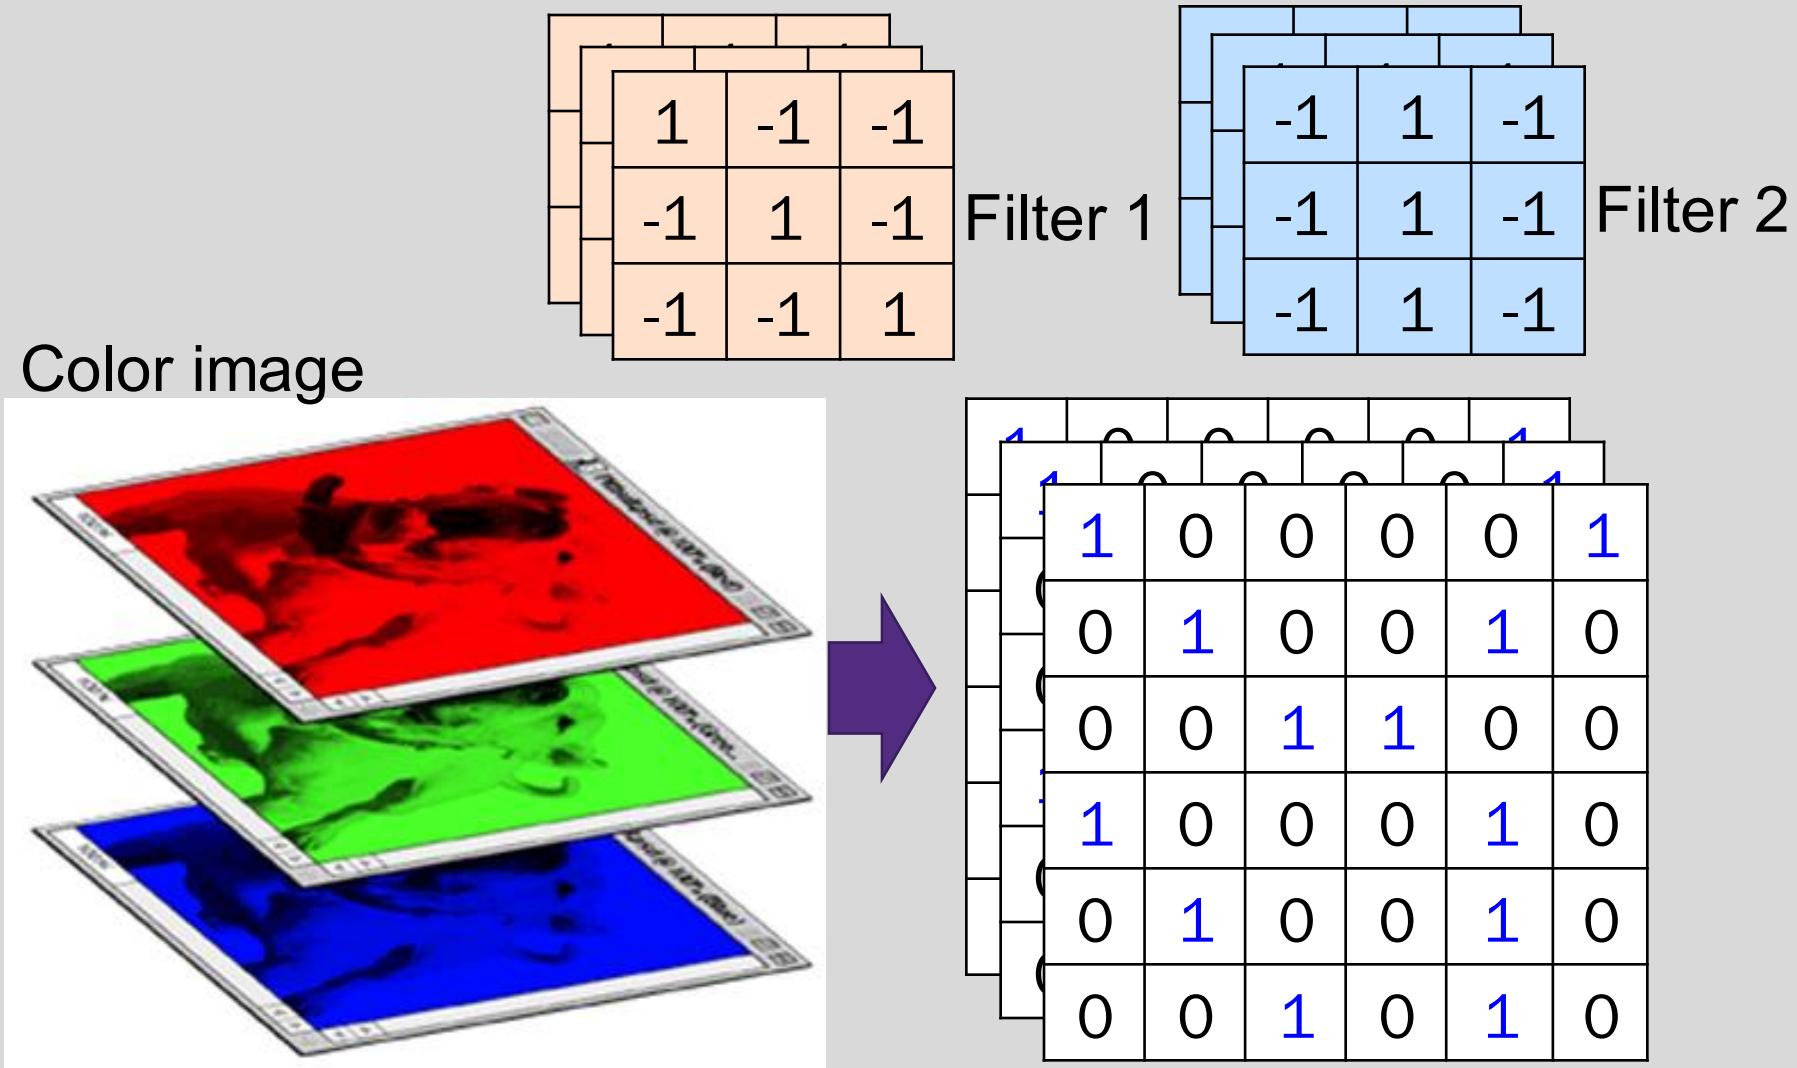
stride=1



6 x 6 image



# Color image: RGB 3 channels



# More efficiency with pooling

Subsampling pixels will not change the overall scene

bird



Subsampling

bird



We can subsample the pixels to make image smaller



fewer parameters to characterize the image

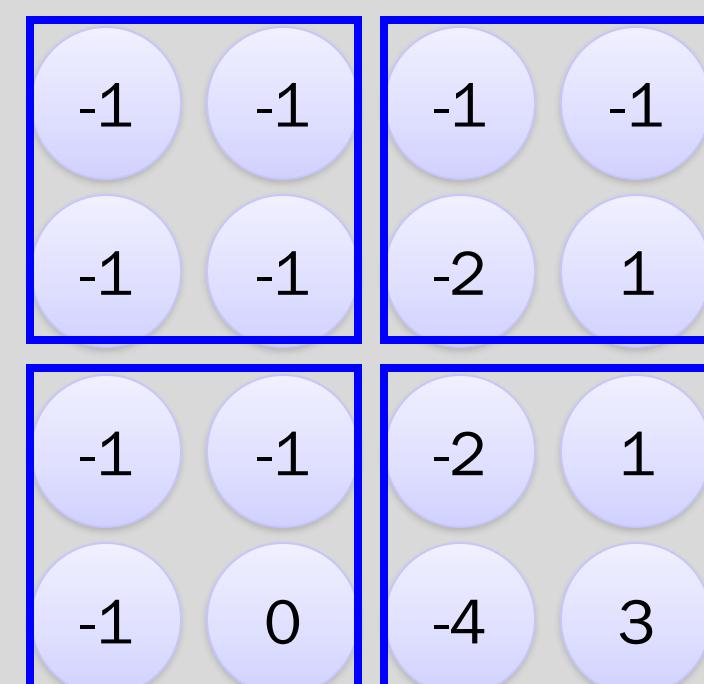
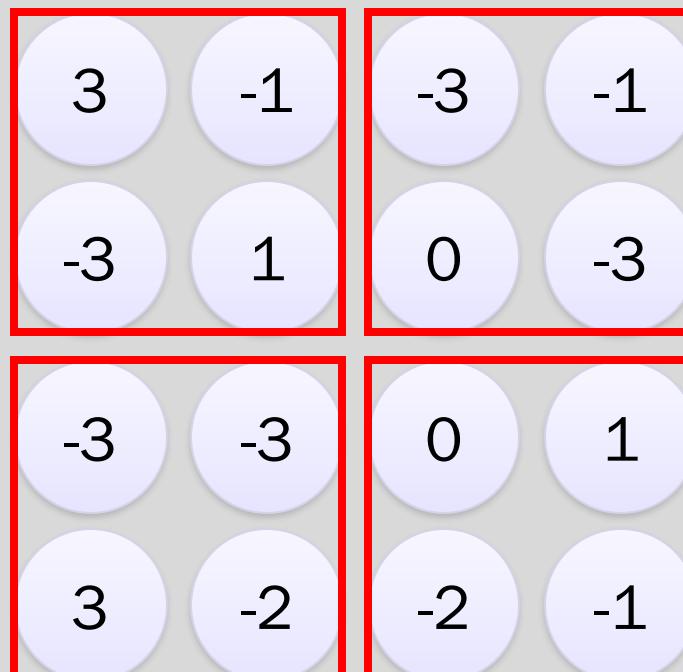
# Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2



# Convolution and Pooling

Original Image

1	1	1	1	1	1
1	1	1	1	1	1
0	0	1	1	0	0
0	0	1	1	0	0
0	0	1	1	0	0
0	0	1	1	0	0

Feature Map

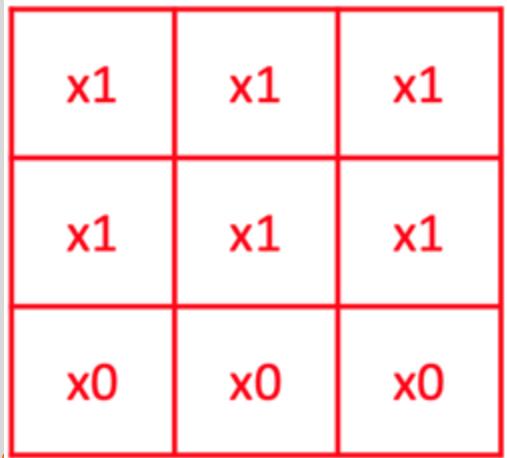

Feature Map

6	6	6	6
4	5	5	4
2	4	4	2
2	4	4	2

Max  
Pooling

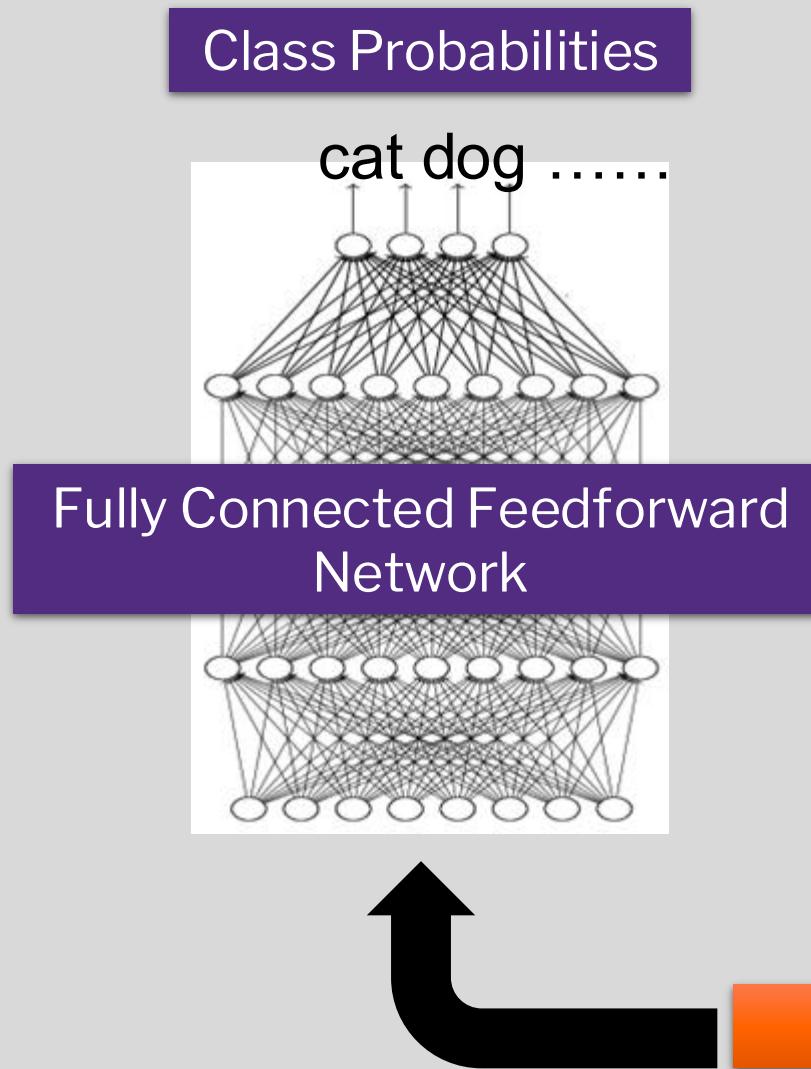

Average  
Pooling


Sum  
Pooling

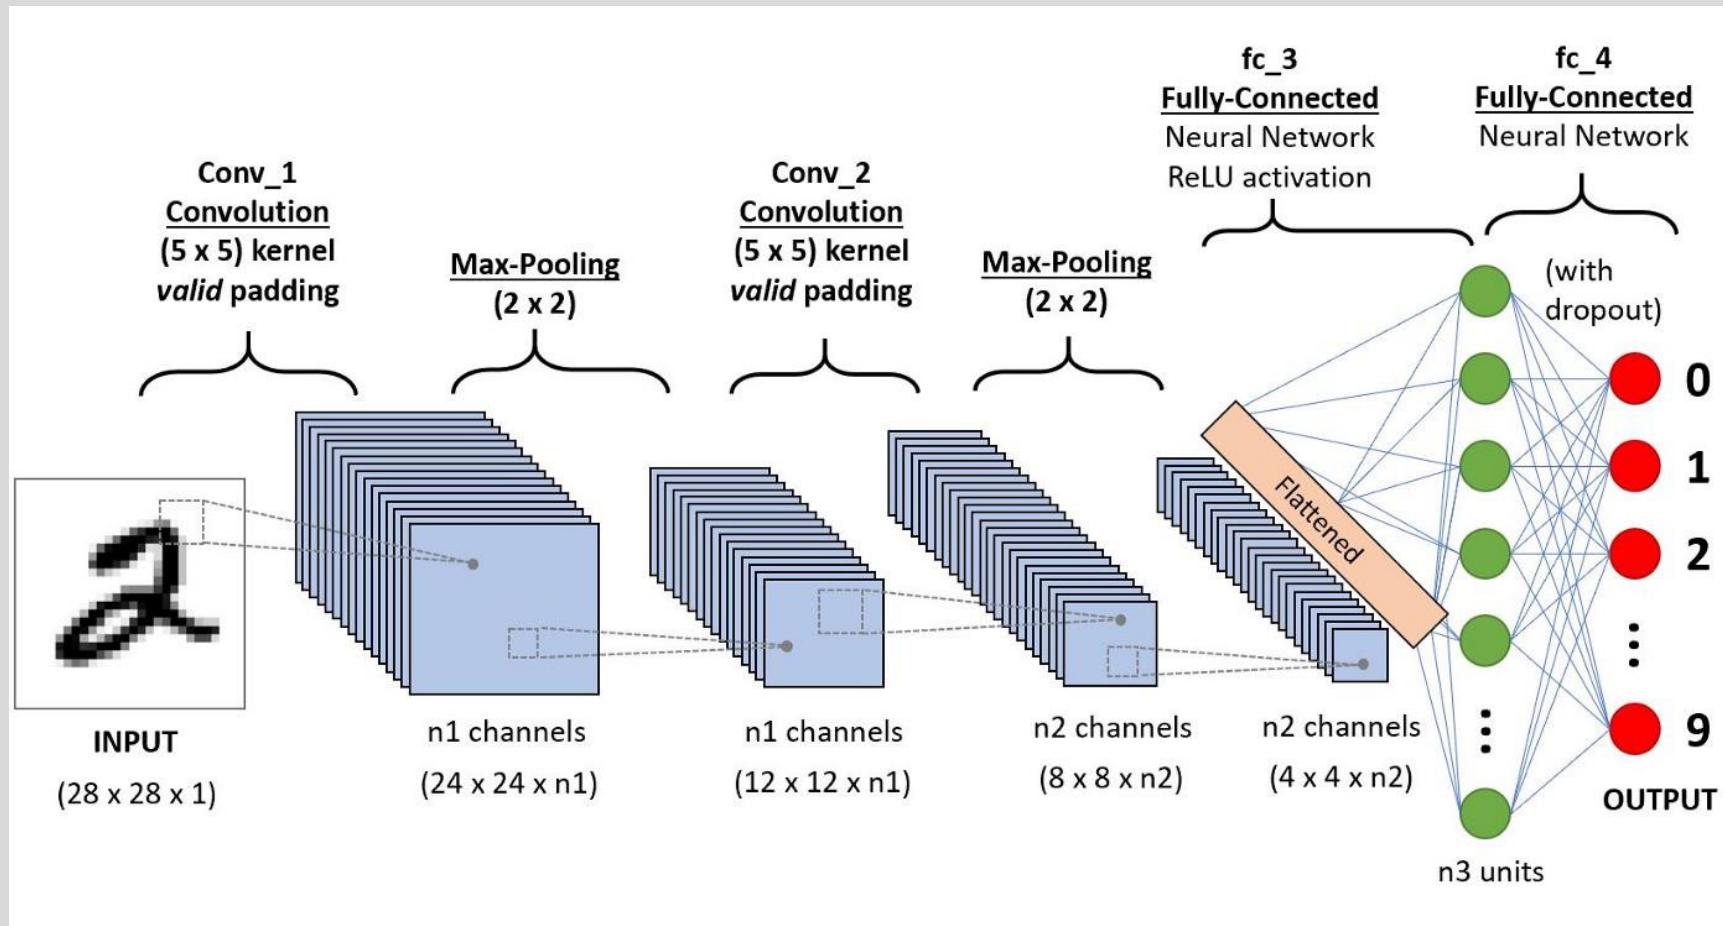



Max pooling is common, but average and sum pooling are also used

# The whole CNN



# Deep CNN Architecture

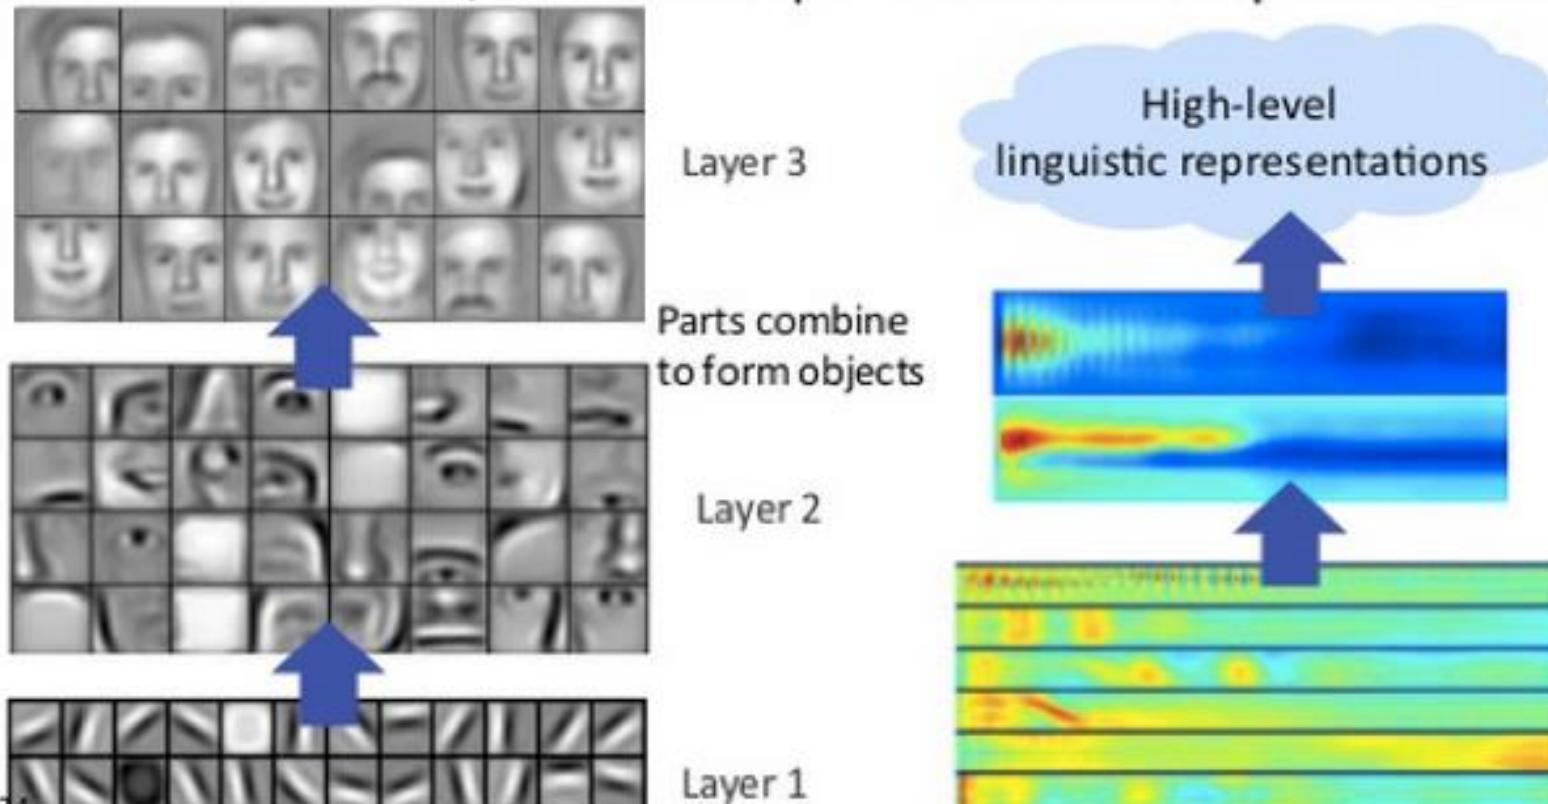


- Multiple convolution layers stacked in sequence
- Each layer can have an arbitrary number of convolution kernels
- Pooling layers for down sampling

<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

# Feature hierarchy

Successive model layers learn deeper intermediate representations



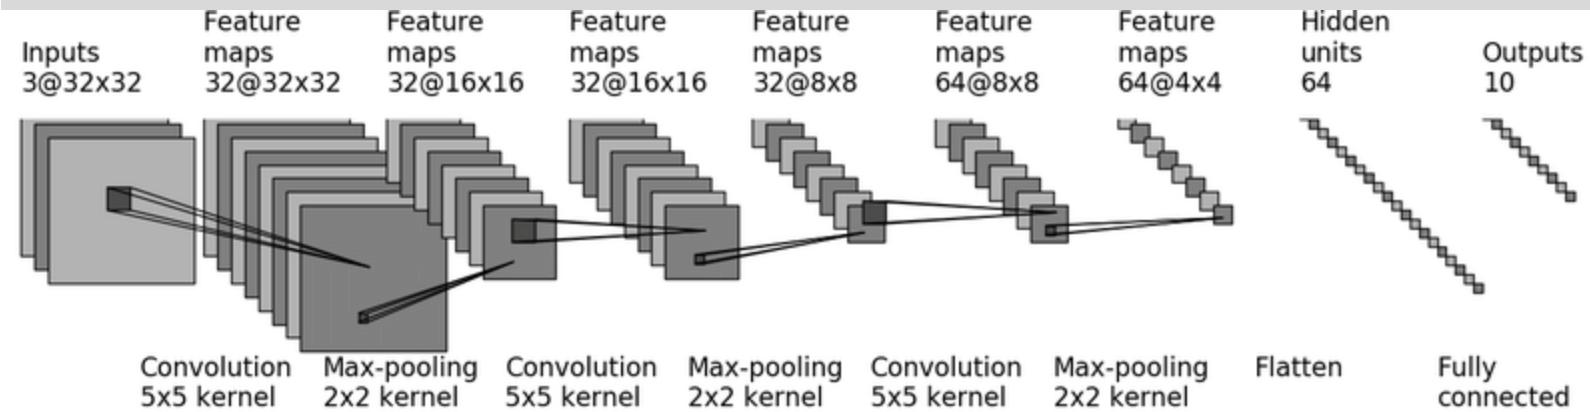
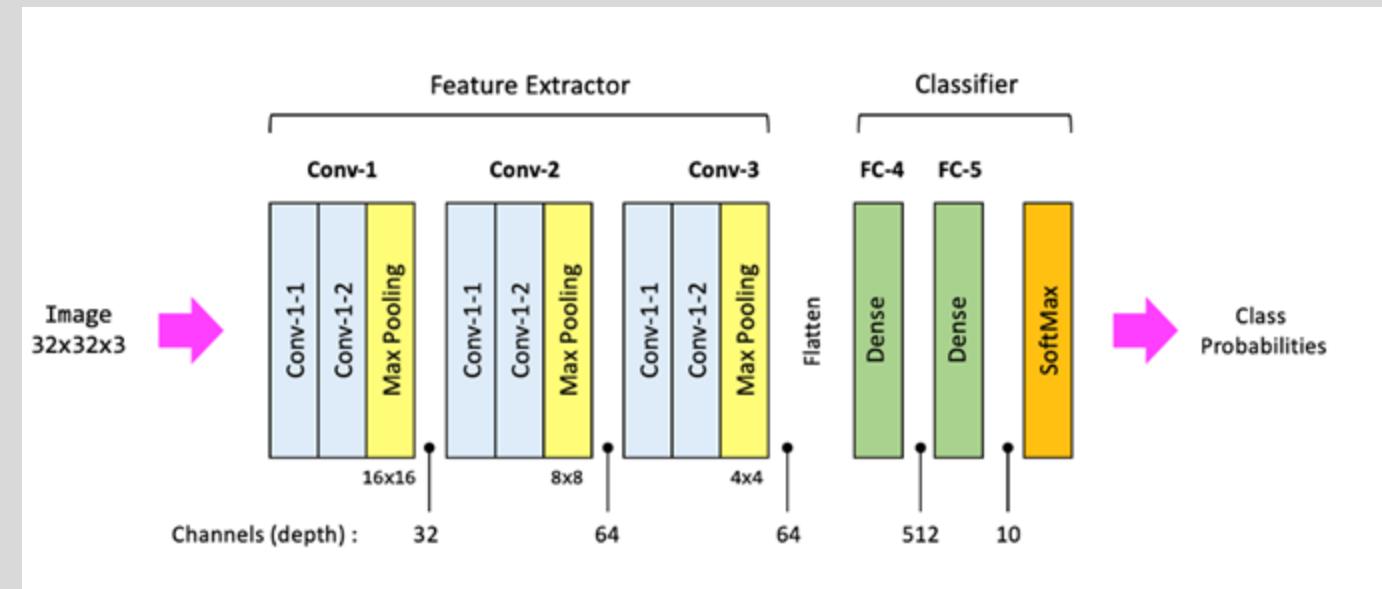
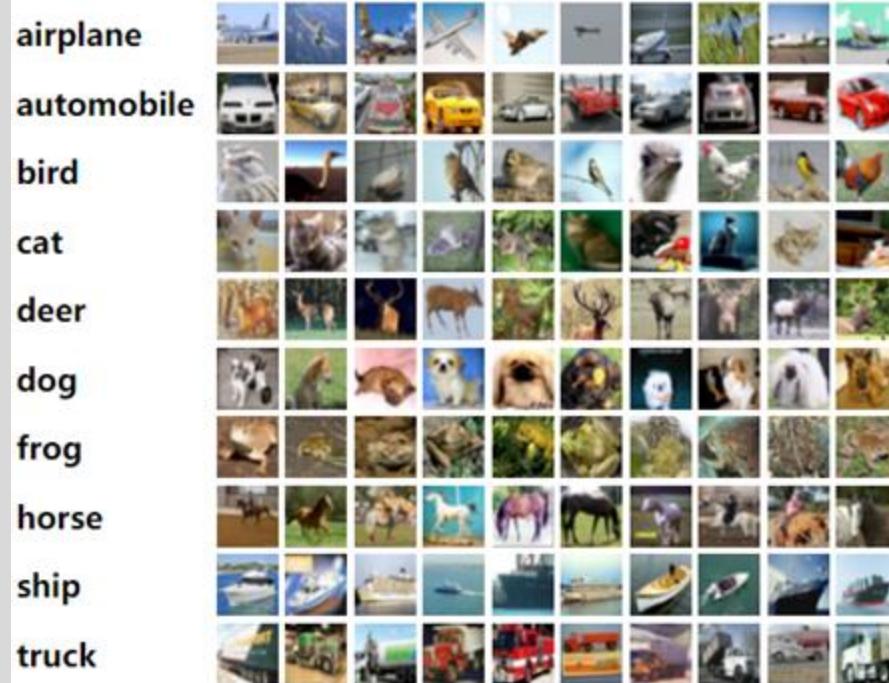
Prior: underlying factors & concepts compactly expressed w/ multiple levels of abstraction

- Early layers tend to learn simple features such as edges
- Feature learning is spatially invariant
- Later layers learn features composed of simpler features to recognize more complex patterns
- Final layers are effectively “object models”

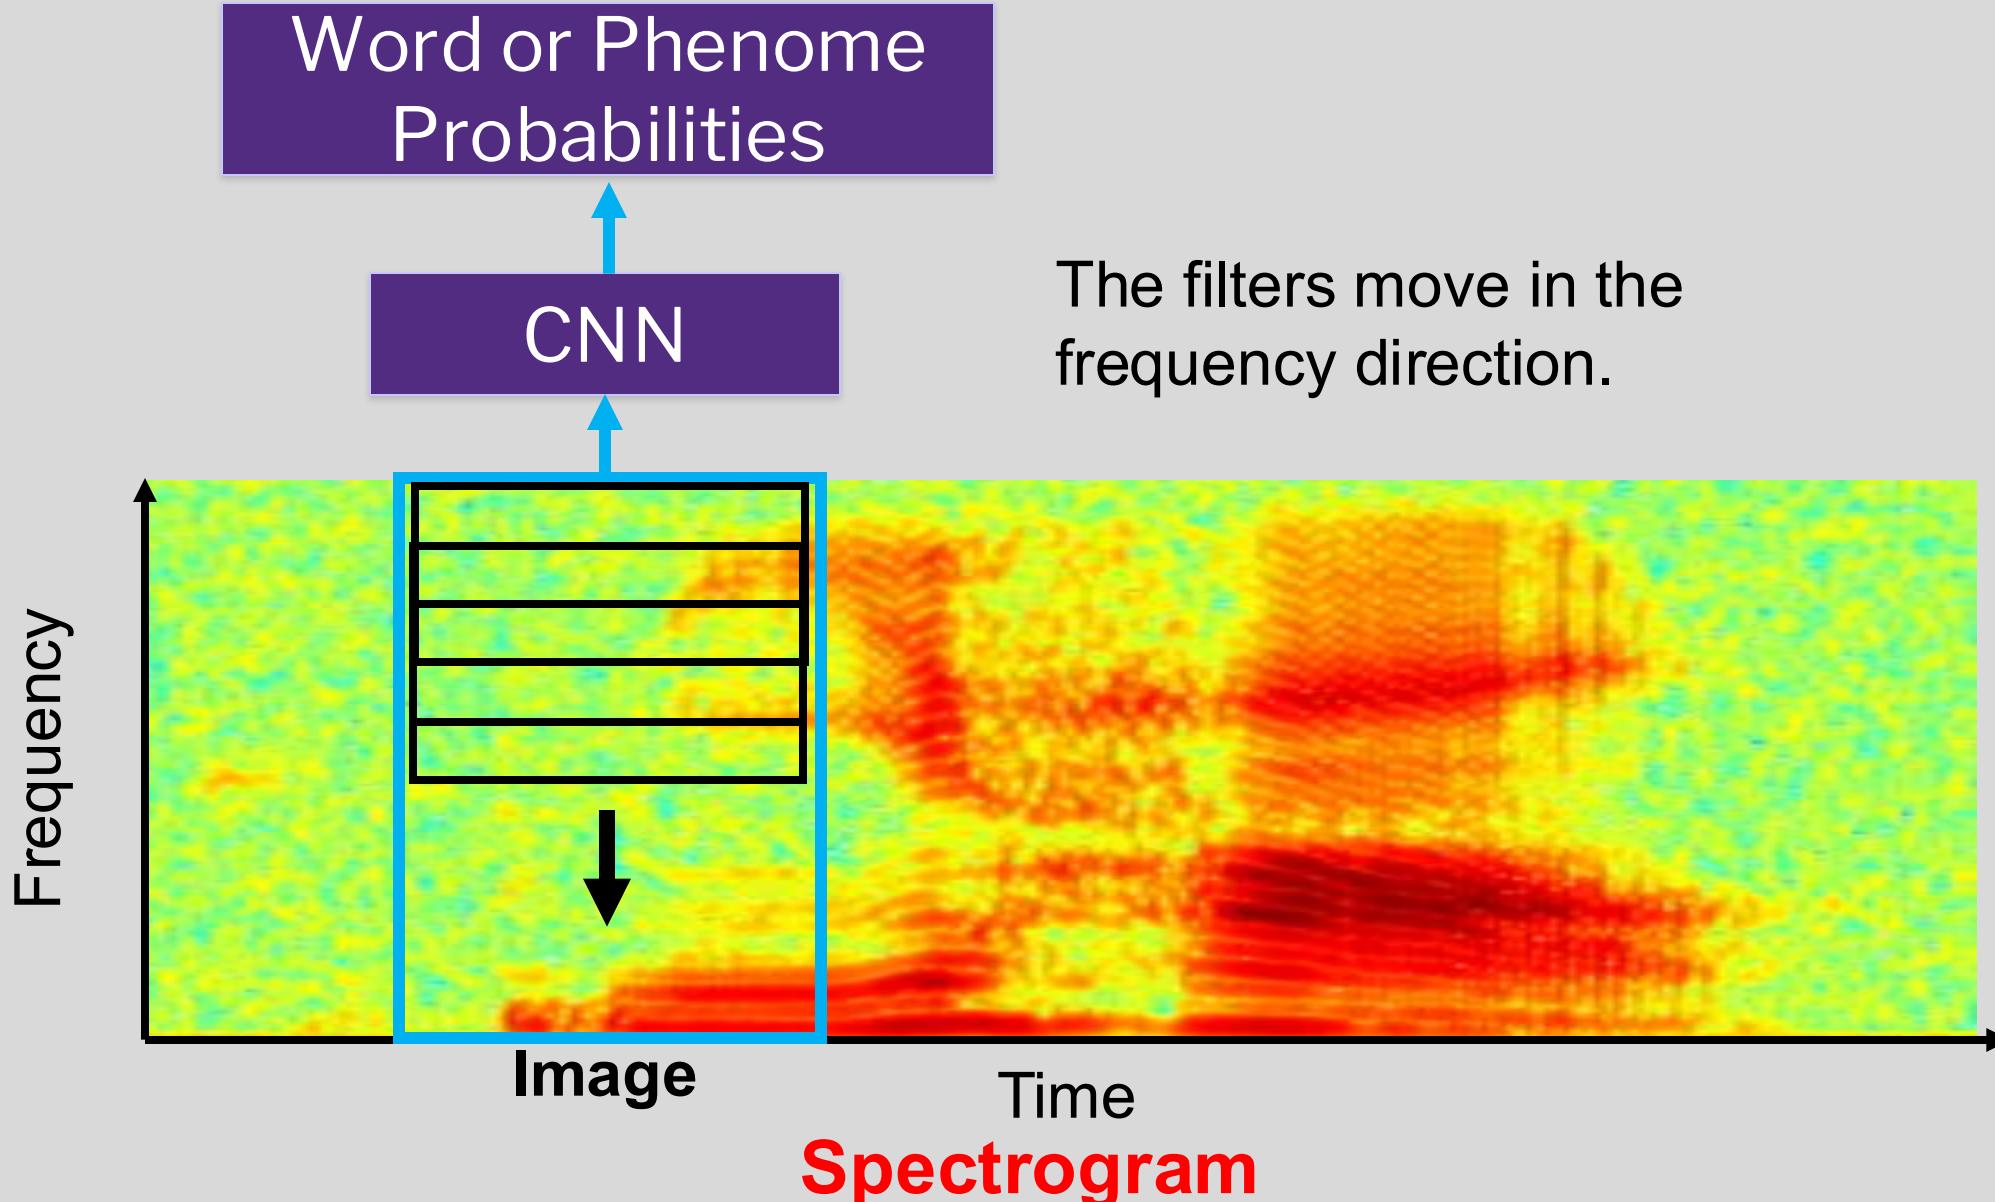
[https://web.stanford.edu/class/archive/cs/cs224n/cs224n.1162/handouts/CS224N\\_DeepNLP\\_Week7\\_lecture1.pdf](https://web.stanford.edu/class/archive/cs/cs224n/cs224n.1162/handouts/CS224N_DeepNLP_Week7_lecture1.pdf)

# CNN for Image Classification

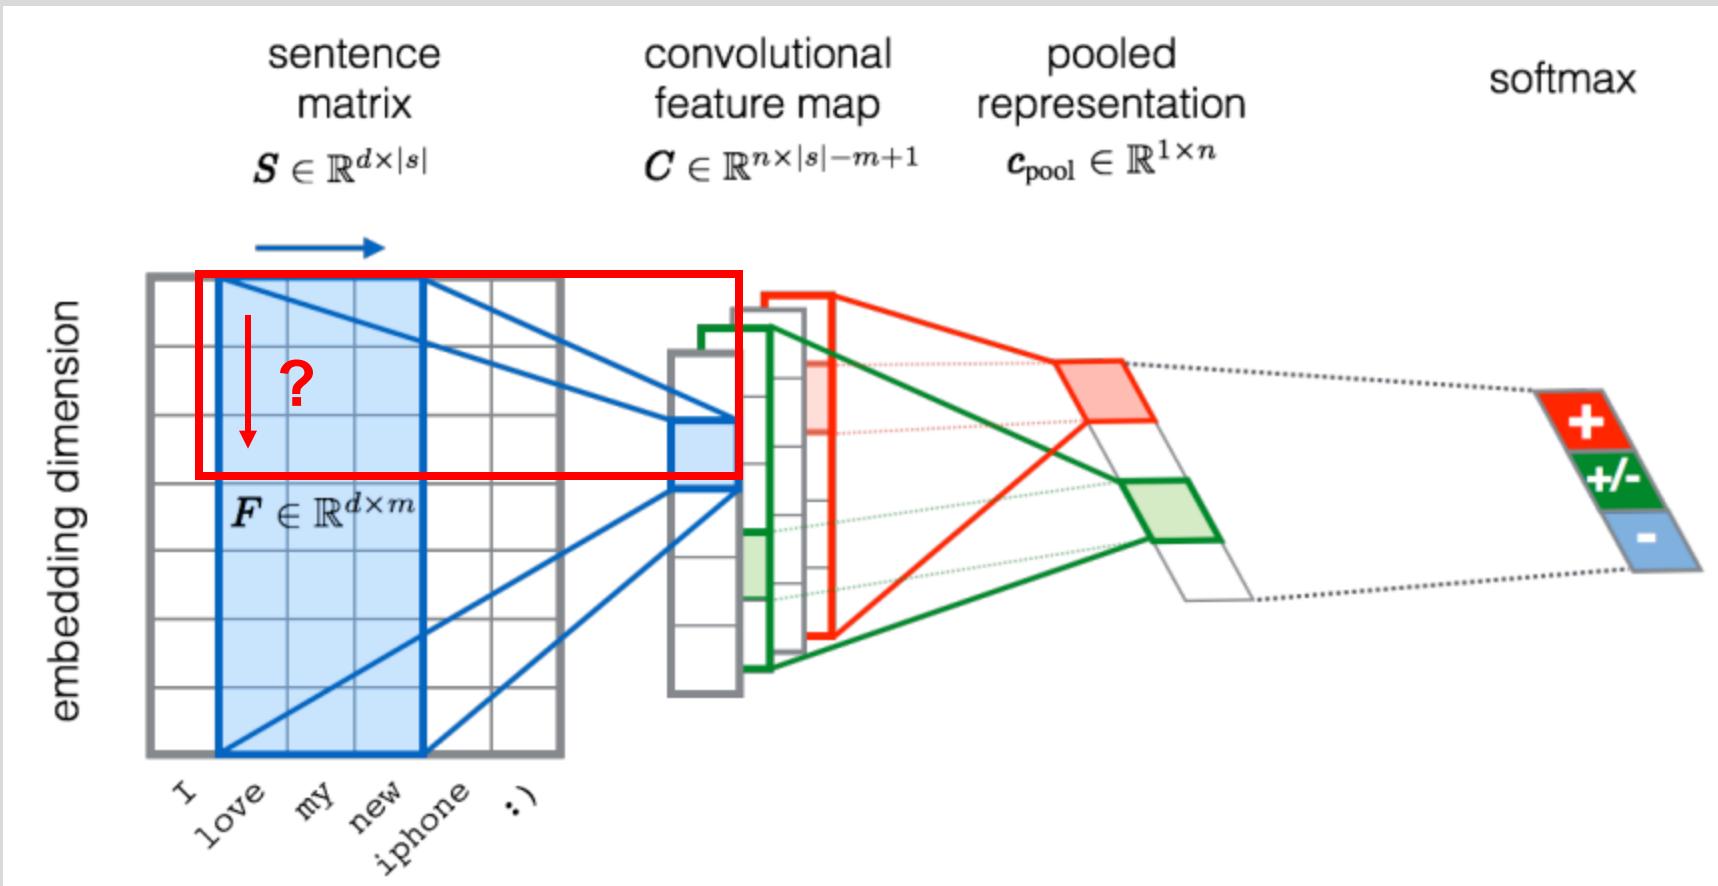
Example using the CIFAR10 dataset, 32 x 32



# CNN for speech recognition



# CNN in text classification



Source of image:  
<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.703.6858&rep=rep1&type=pdf>



# Recurrent Neural Networks (RNN)

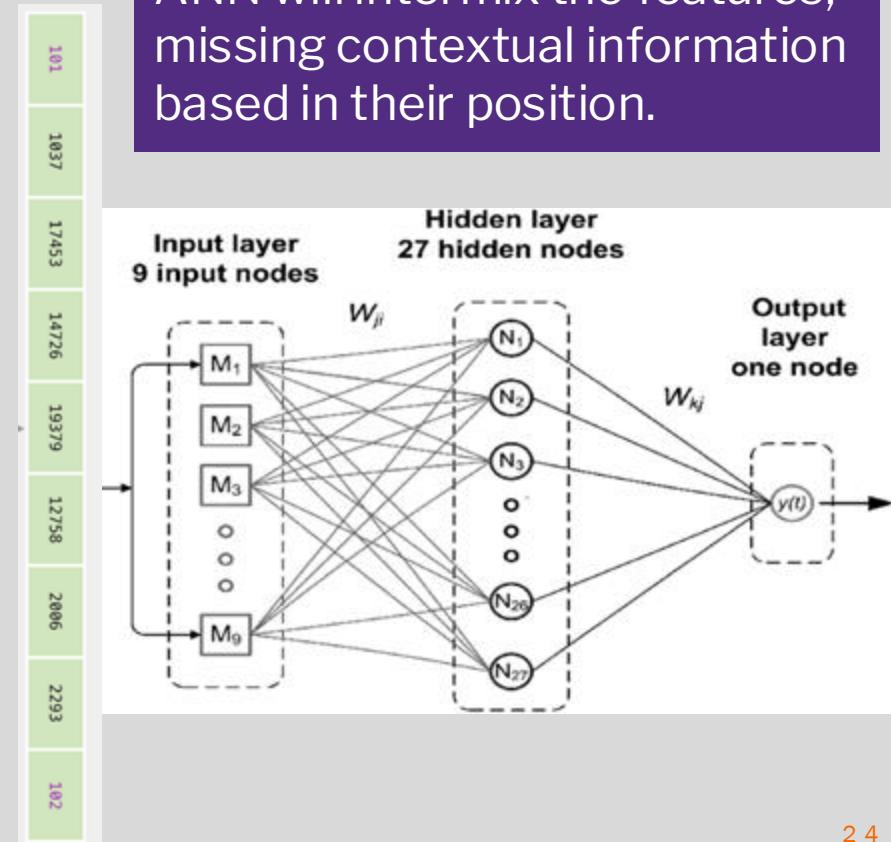


# ANN vs RNN

ANNs do not work well for data where the relative order of inputs is informative and dynamic such as text (sequence of words) or videos (sequence of images)



ANN will intermix the features, missing contextual information based in their position.



# Sequence Learning Applications

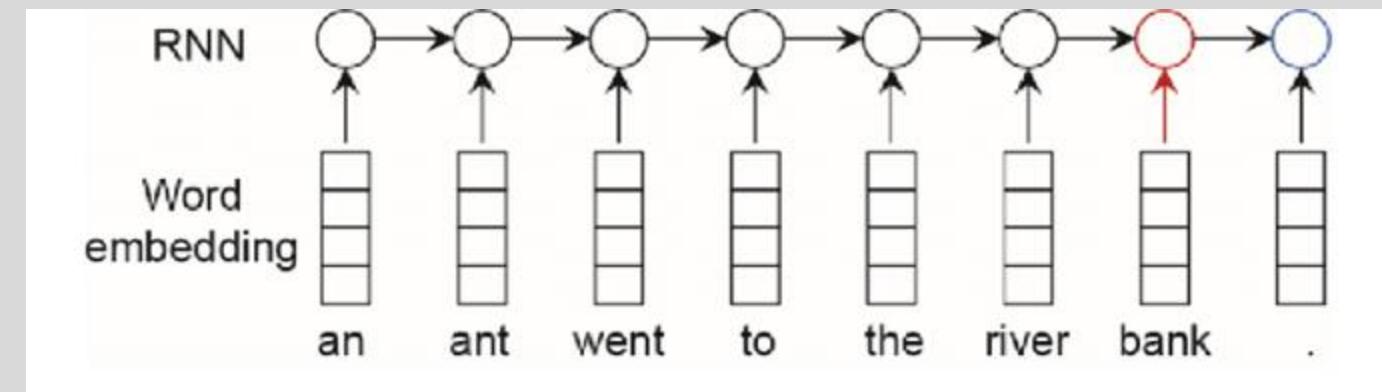
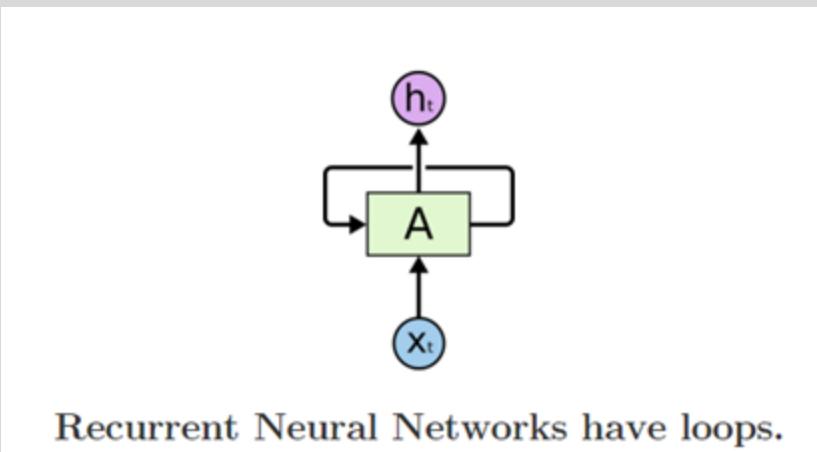
RNNs can be applied to various type of sequential data to learn the temporal patterns that are informative to a given task.

Many applications

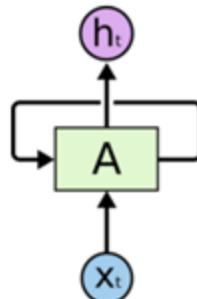
- Time-series data (e.g., stock price) → prediction
- Raw sensor data (e.g., signal, voice, handwriting) → sequence of class labels
- Text document → single label (document level), multiple labels (document partition labels)
- Image and video → Text description (e.g., captions, scene interpretation)

# Recurrent Neural Network (RNN)

- RNN units (or cells) add loops to retain prior information
- Inputs are processed in sequential order



# Recurrent Neural Networks



Recurrent Neural Networks have loops.

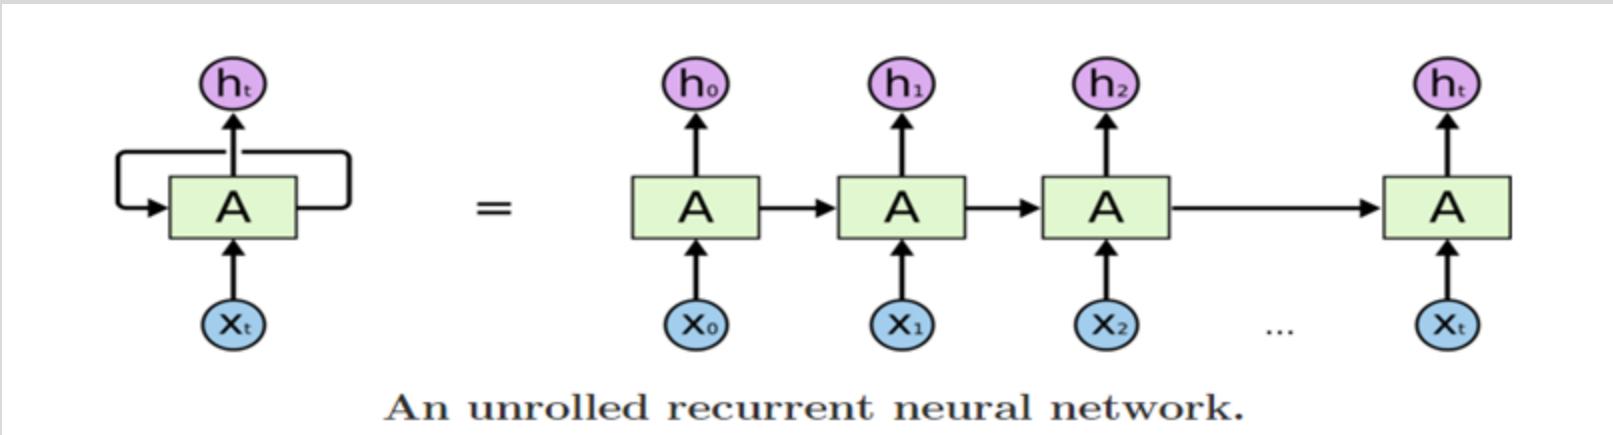
In the above diagram, a chunk of neural network,  $\mathbf{A} = f_{\mathbf{W}}$ , looks at some input  $\mathbf{x}_t$  and outputs a value  $\mathbf{h}_t$ . A loop allows information to be passed from one step of the network to the next.

$$\text{new state } \boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$

function with parameter  $W$

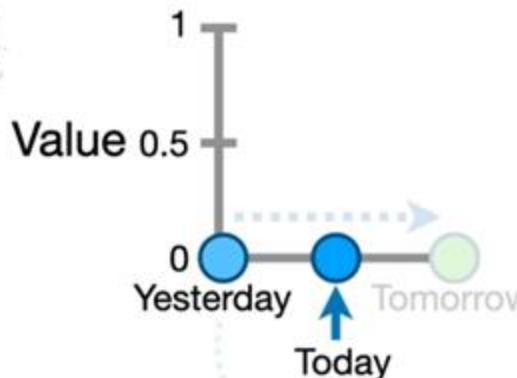
Input vector at some time step

# “Unrolling” Recurrent Neural Networks



A recurrent neural network can be thought of as multiple copies of the same network, each passing a message to a successor. The diagram above shows what happens if we **unroll the loop**.

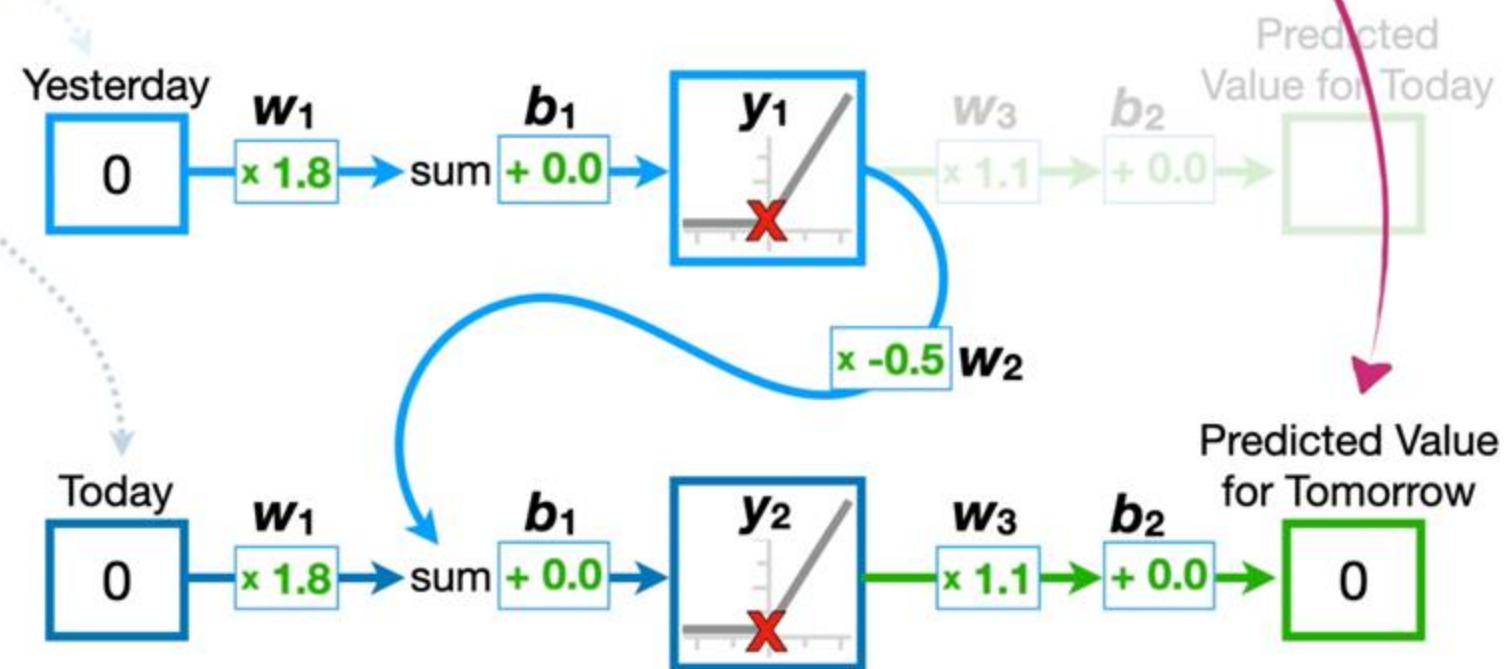
# “Message” passing in RNNs



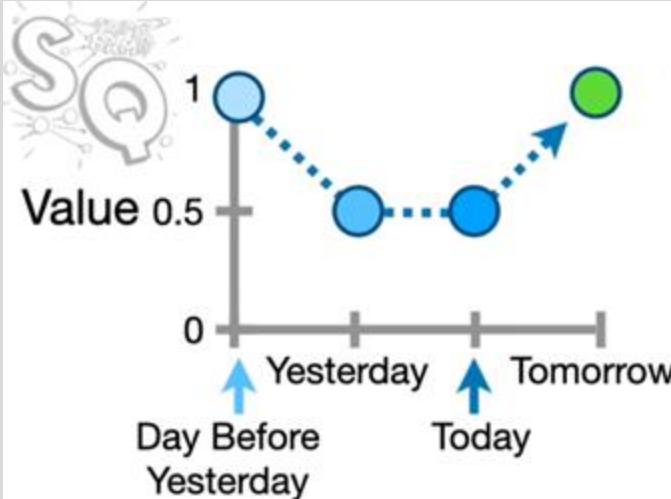
...and that gives us a predicted value for tomorrow, 0...

$$y_2 \times w_3 + b_2 = \text{Prediction for Tomorrow}$$

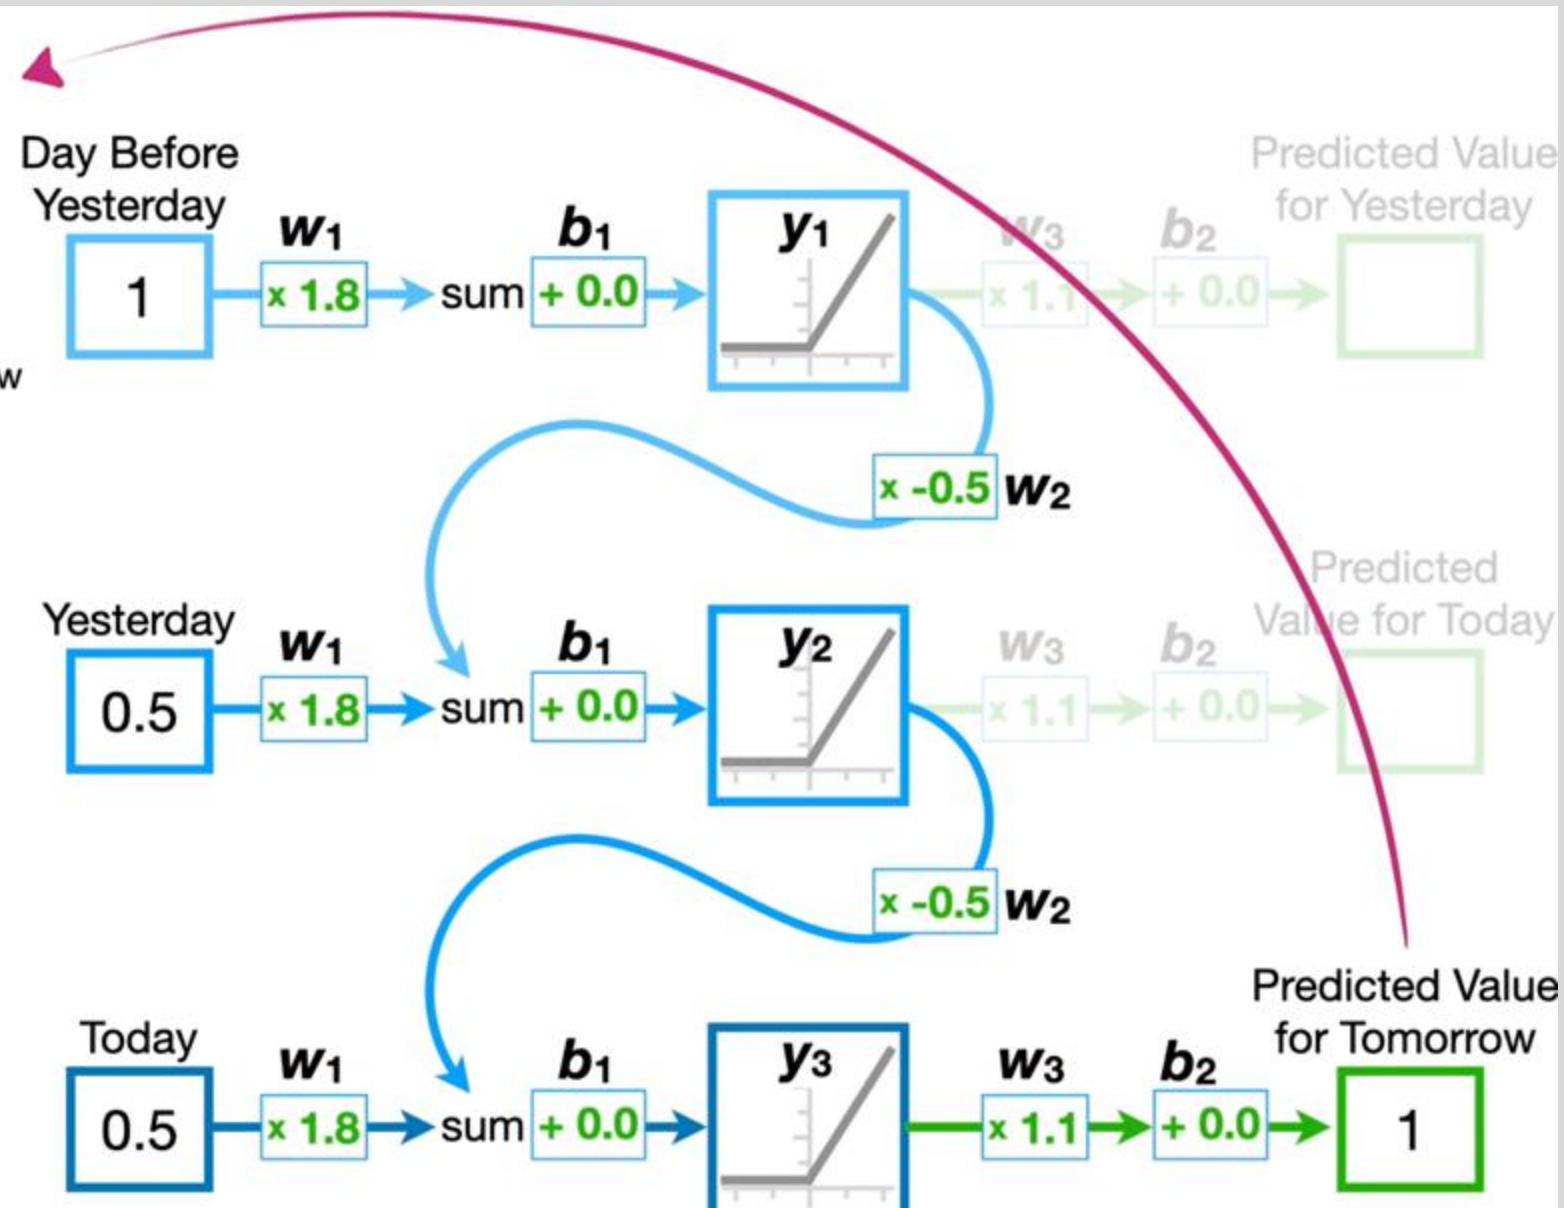
$$0 \times 1.1 + 0.0 = 0$$



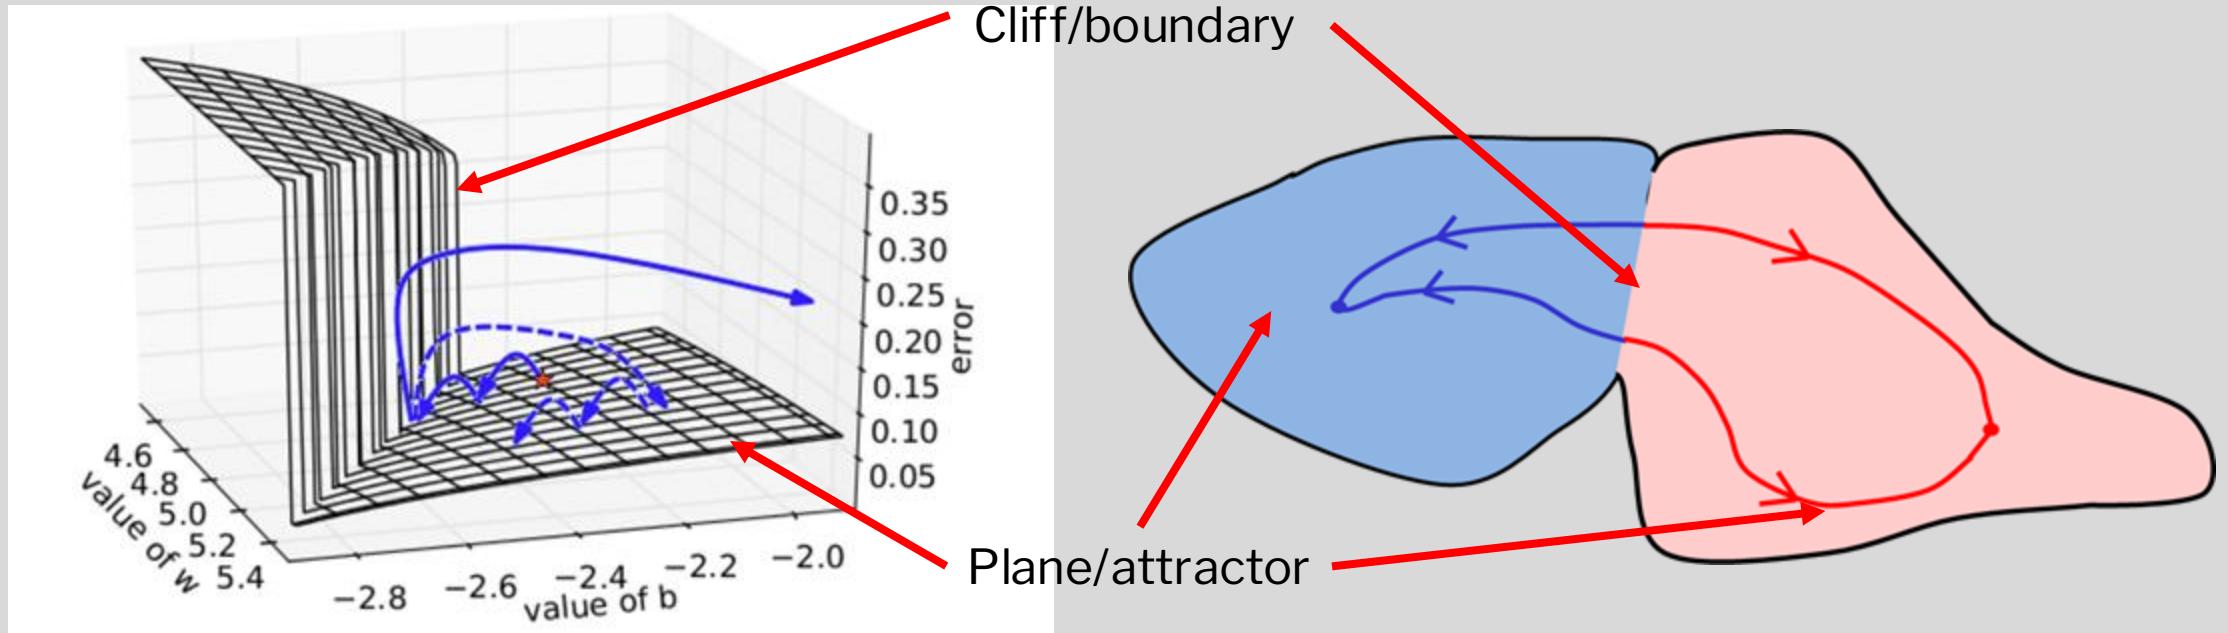
# “Message” passing in RNNs



And when we do the math, the last output gives us the prediction for **tomorrow**.



# Exploding and Vanishing Gradients



**Exploding:** If we start almost exactly on the boundary (cliff), tiny changes can make a huge difference.

**Vanishing:** If we start a trajectory within an attractor (plane, flat surface), small changes in where we start make no difference to where we end up.

Both cases hinder the learning process.

# Exploding gradient problem

Large numbers of  
 $W_2$  can rapidly  
explode.



And that means the first  
input value is amplified  
**16 times** before it gets  
to the final copy of the  
network.

$$\text{Input}_1 \times 2 \times 2 \times 2 \times 2$$

$$= \text{Input}_1 \times 2^4 = \text{Input}_1 \times 16$$

$$= \text{Input}_1 \times w_2^{\text{Num. Unroll}}$$



# Vanishing gradient problem

Small values of W2  
can rapidly vanish



And that means the first  
input value is amplified  
**16 times** before it gets  
to the final copy of the  
network.

|

$$\begin{aligned} Input_1 \times (1/2)^4 \\ = \frac{Input_1}{16} \end{aligned}$$



# Networks with Memory

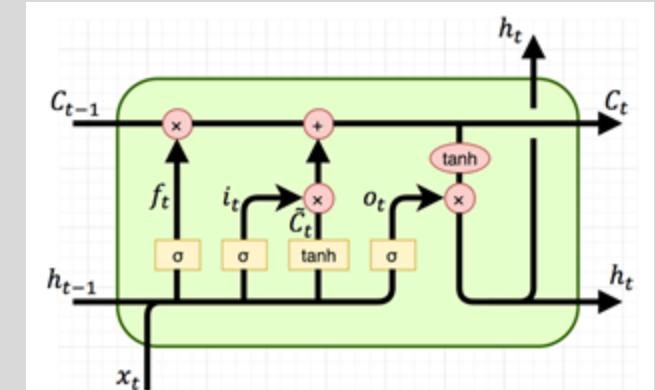
Standard (aka vanilla) RNN operates in a “multiplicative” way leading to vanishing/exploding gradients

Two recurrent cell designs have proposed and widely adopted:

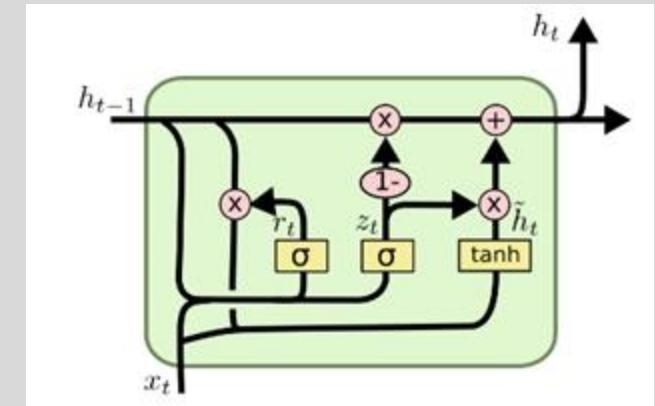
- Long Short-Term Memory (LSTM) (Hochreiter and Schmidhuber, 1997)
- Gated Recurrent Unit (GRU) (Cho et al. 2014)

Both designs process information in an “additive” way with gates to control information flow.

- Sigmoid gate outputs numbers between 0 and 1, describing how much of each component should be let through.



Standard LSTM Cell

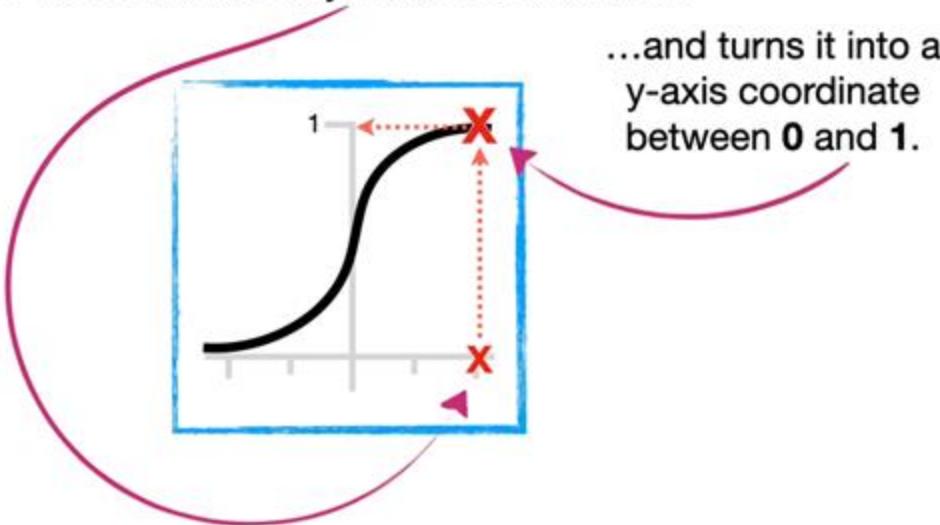


GRU Cell

# Activation functions

In a nutshell, the **Sigmoid Activation Function** takes any x-axis coordinate...

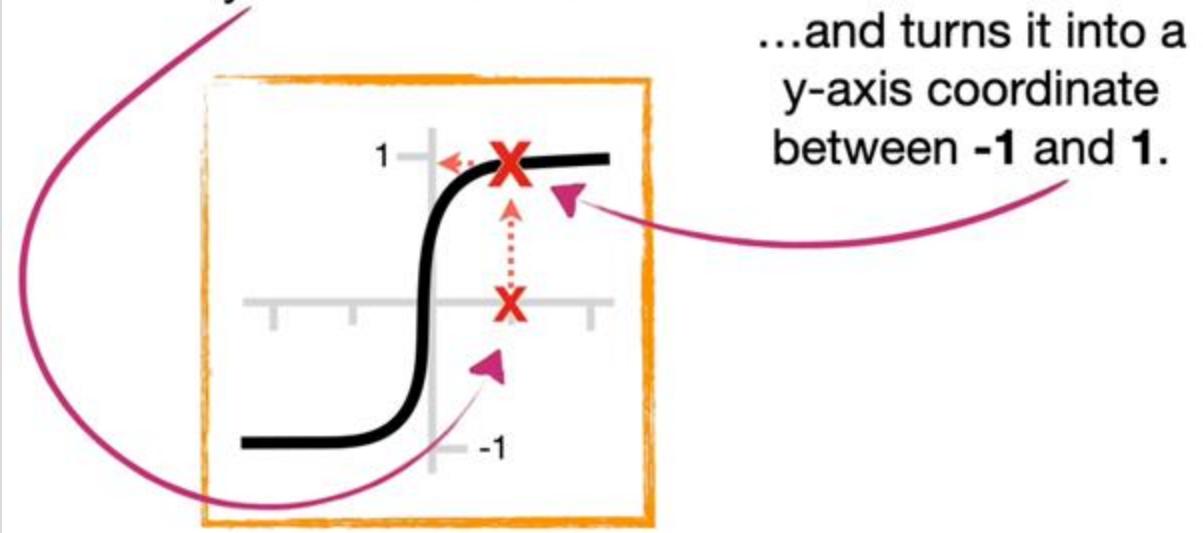
...and turns it into a y-axis coordinate between **0** and **1**.



Useful for gates, where we want no (0) or all (1) information to flow

**Tangent Activation Function** takes any x-axis coordinate...

...and turns it into a y-axis coordinate between **-1** and **1**.



Useful for cell outputs where we want to avoid vanishing gradients (ReLU can also be used)

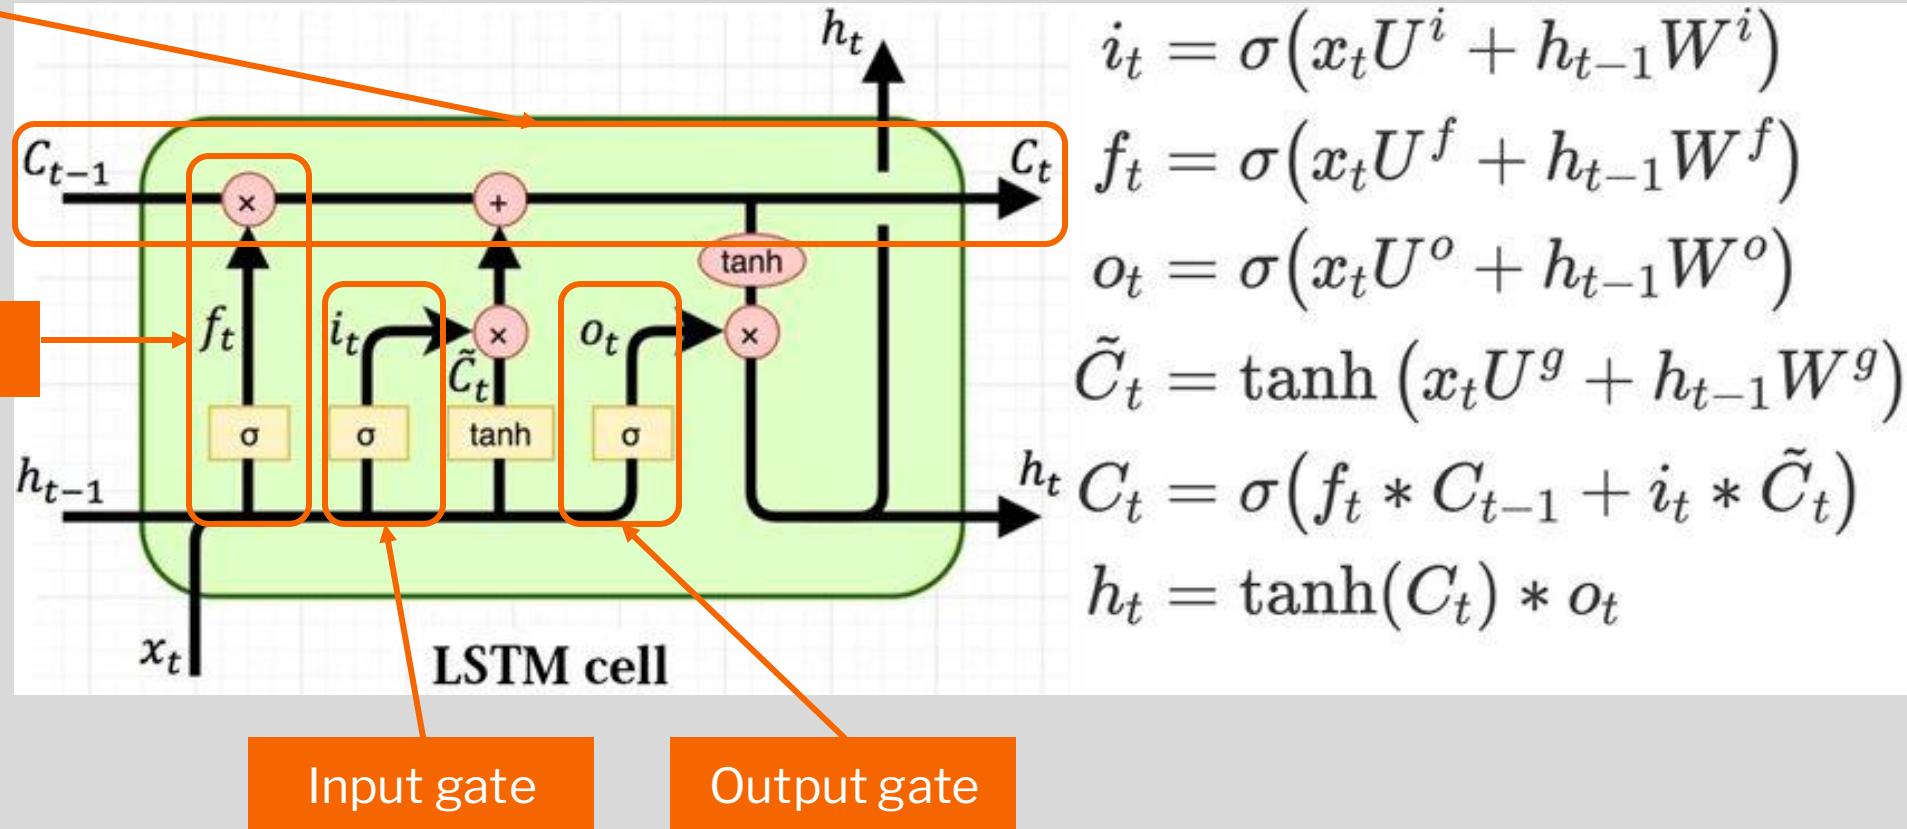
# LSTM

Cell state  
(long term  
memory)

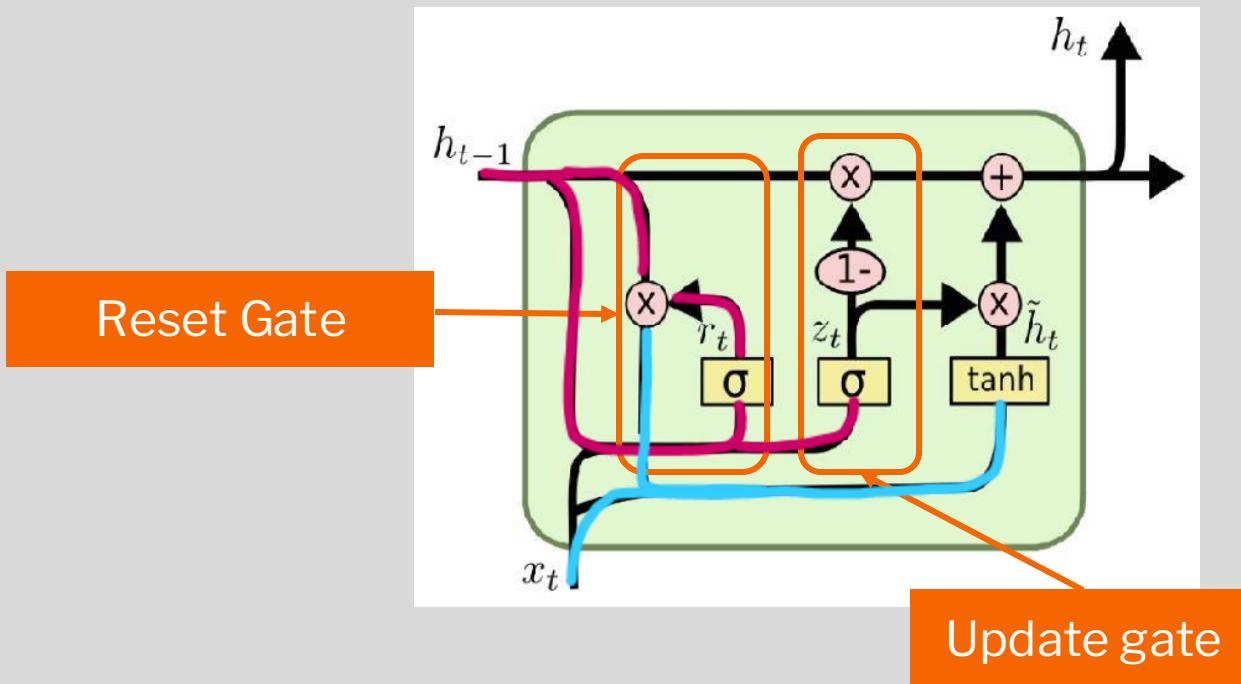
Forget gate

Input gate

Output gate



<https://www.youtube.com/watch?v=YCzL96nL7j0>



$$z_t = \sigma (W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

<https://www.youtube.com/watch?v=YCzL96nL7j0>

# Exploding & vanishing gradients revisited

- The exploding / vanishing gradient problem also occurs in other networks types when many layers are used (e.g., several CNN layers)
- Mitigation strategies include:
  - ReLU (and its variants) activation function
  - Weight initialization via Xavier, He, Kaiming methods
  - Gradient clipping
  - Batch normalization



# Questions?