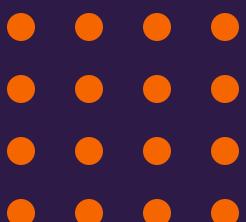




# Introduction to Graph Neural Networks

**Aaron J. Masino, PhD**  
Associate Professor, School of Computing



# Outline

- Graph terminology
- Graph neural networks



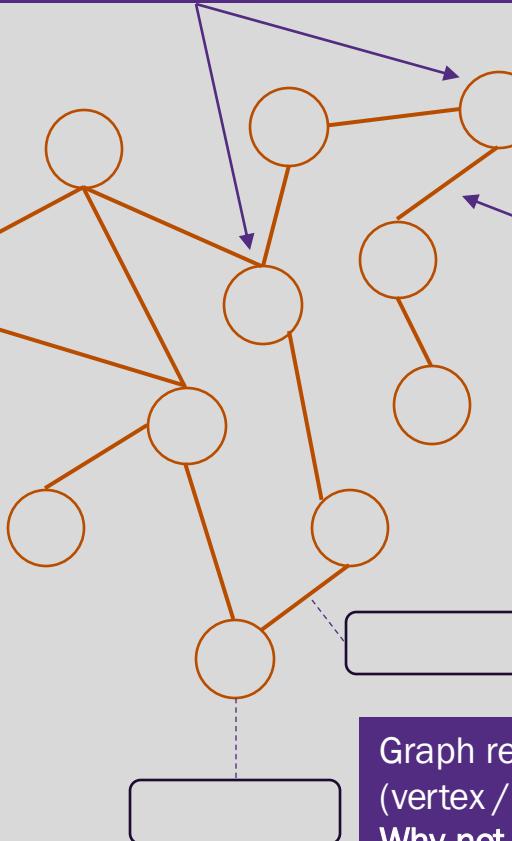
# Introduction to graphs



# Terminology

Vertices (nodes) usually represent entities (e.g., a patient, symptom, disease, cell type, atom/molecule, gene)

Neighbors – two nodes connected by an edge  
Neighborhood – all neighbors of a node



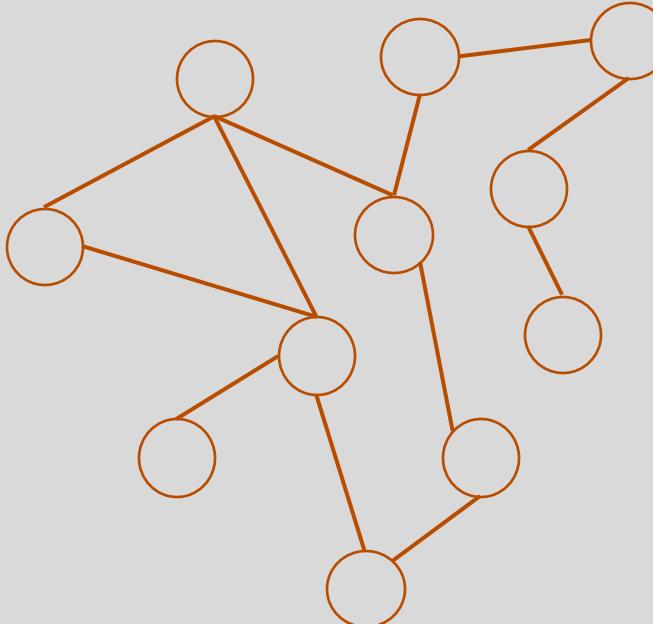
Edges usually represent relationships between nodes (e.g., “has symptom”, “causes disease”, “bonds with”, “regulates expression of”).

Graph representations may also include *attributes* (vertex / edge properties, e.g., text description).  
**Why not just represent these as additional vertices?**

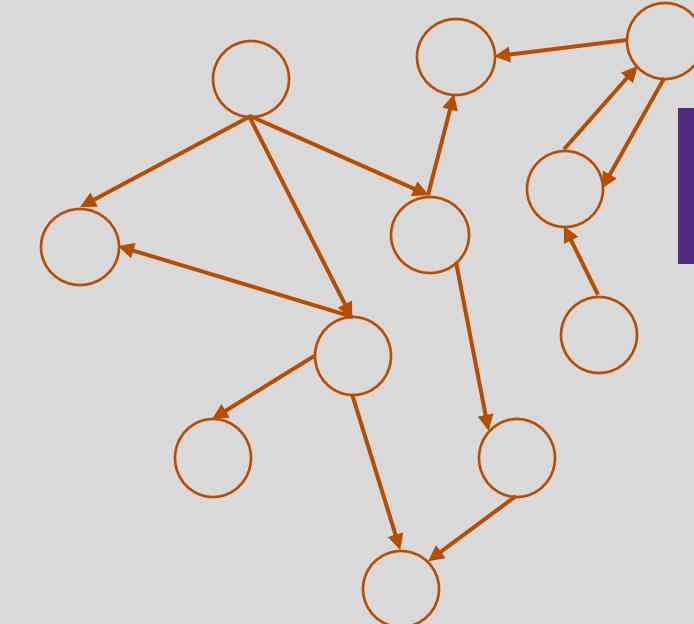
- Tends to lead to very dense networks for properties that are common to most nodes.
- Such properties usually don't describe relations between nodes
- Computational efficiency

# Terminology

**Undirected Graph** – edges have no direction and define relations that are inherently bidirectional.

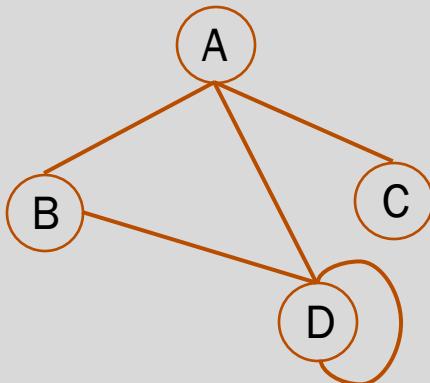


**Directed Graph** – edges have a specified direction and define relations that are inherently unidirectional.

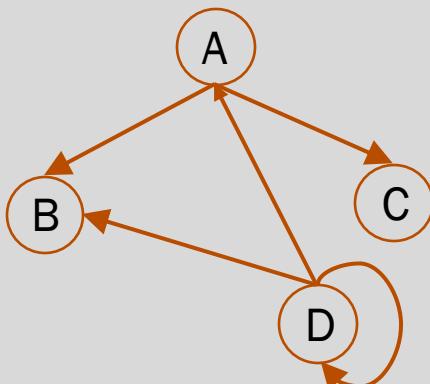


Directed graphs may contain edges in both directions between nodes.

# Adjacency matrix



	A	B	C	D
A	0	1	1	1
B	1	0	0	1
C	1	0	0	0
D	1	1	0	1

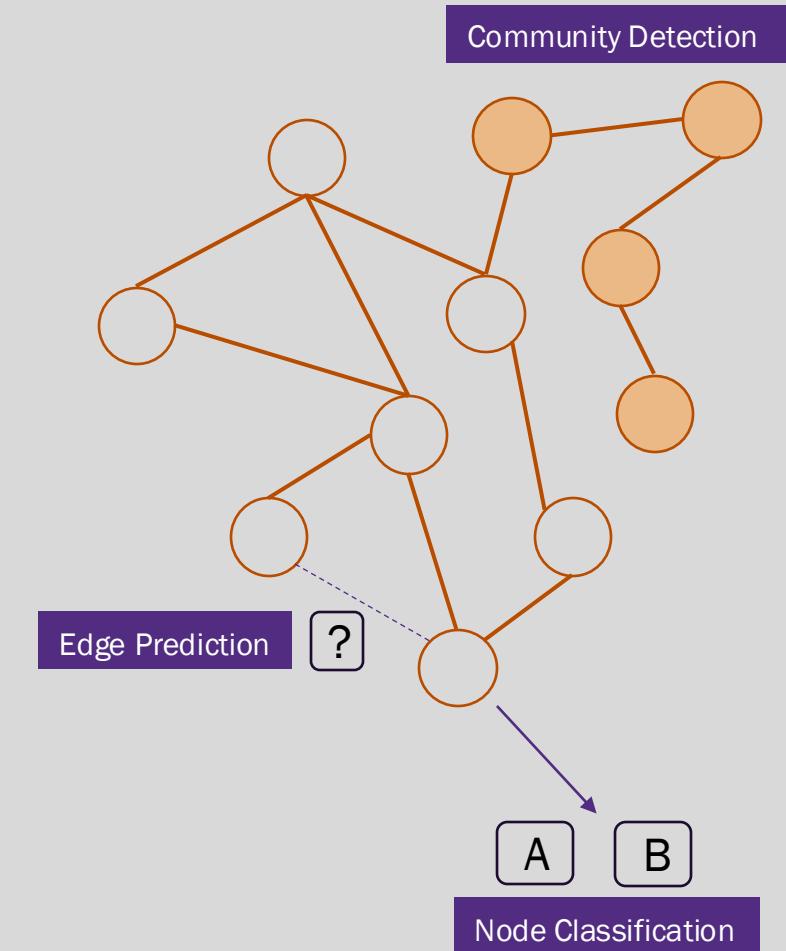


	A	B	C	D
A	0	1	1	0
B	0	0	0	0
C	0	0	0	0
D	1	1	0	1

- Adjacency matrix,  $A$ , summarizes the edges in a graph where the element  $a_{i,j} \in A$  is the weight of the edge  $e_{i,j} \in E$ . For unconnected vertices,  $v_i, v_j, a_{i,j} = 0$
- Adjacency matrix is symmetric for undirected graphs and usually asymmetric for directed graphs
- For graphs without self loops,  $a_{i,i} = 0$ , i.e., the diagonal elements are all 0

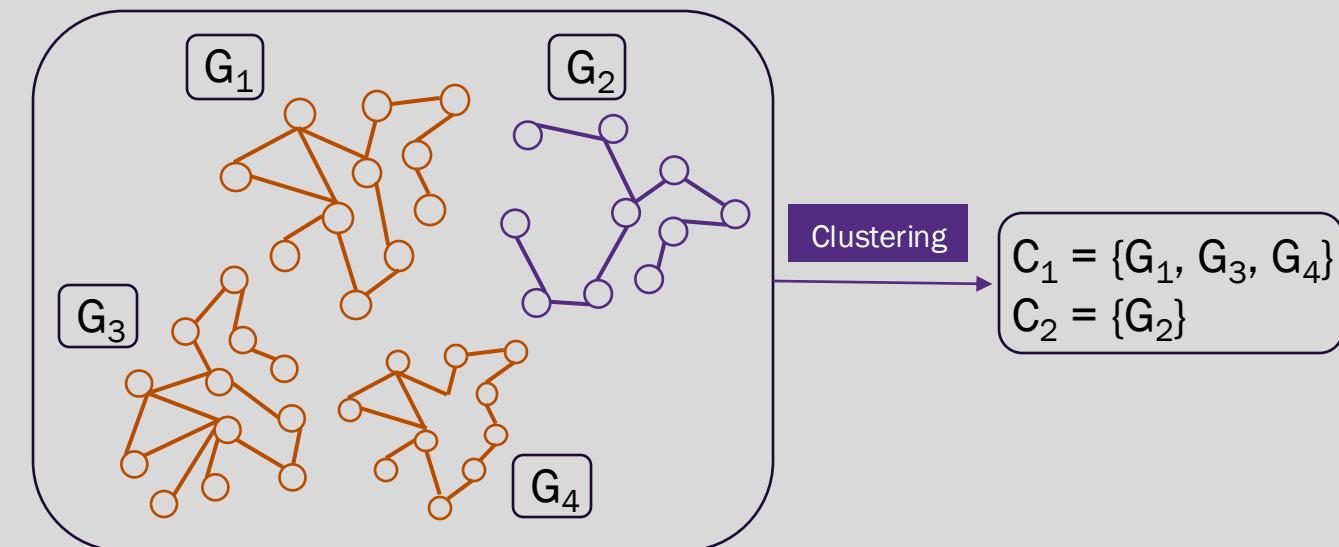
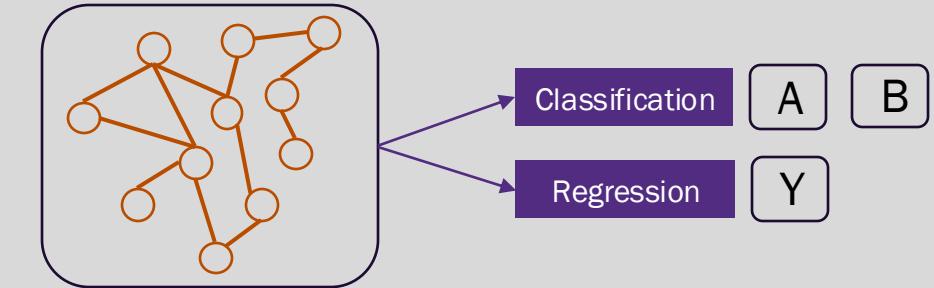
# Intra-graph machine learning tasks

- Node classification – given information about the node (may include neighbors, relationships, attributes), infer class membership (usually an attribute)
- Edge prediction – given information about two nodes, determine if a given relation exists between them
- Community detection (node clustering) – detection of “meaningful” node groups



# Whole graph and inter-graph machine learning tasks

- Whole graph classification and regression – given one entire graph, infer class membership or estimate (regress) a related numerical property
- Inter graph tasks – given multiple graphs estimate the similarity between graphs and/or identify clusters



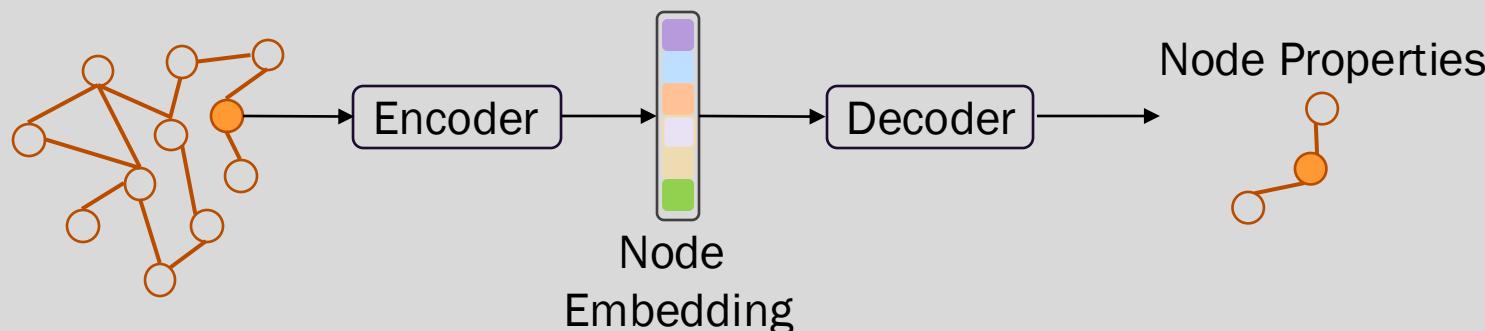


# Graph neural networks



# Graph representation for machine learning

- Traditional approaches utilize engineered features (e.g., node degree, adjacency matrix eigenvalues) to represent nodes, edges, and graphs for analysis
- Graph neural network (GNN) approaches adopt an *encoder – decoder* perspective
  - *encoder* transforms the graph elements (nodes, edges, entire graph) into a latent space
  - *decoder* estimates values of interest (class membership, node attributes, relationships) from the encoded representation

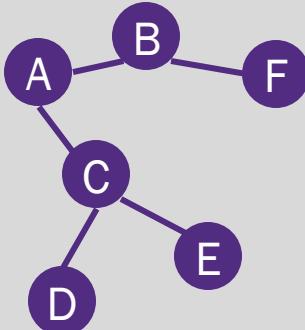




# Random walk methods

- Shallow embedding approaches that attempt to encode random walk statistics
- Goal is for the decoder to approximate the probability distribution of visiting a node  $v$  on a path of length  $T$  starting at node  $u$  on graph  $G$

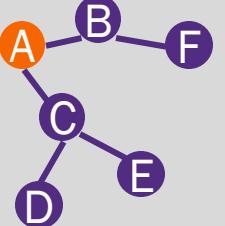
$$P_{G,T}(v|u)$$



## Random Walks

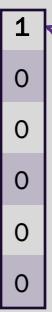
- Assume we have transition states,  $p(u, v | v \in N(u))$  that define the probability of traversing from node  $u$  to  $v$
- Random walks are generated by starting at a randomly selected node,  $u$ , and probabilistically selecting the next node from  $N(u)$  and repeating
- Sample paths from the graph above include
  - ABF
  - ABACE
  - ABADC

# +node2vec



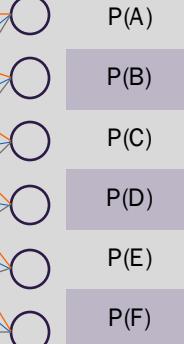
Input Node  
Represented as a  
1-hot encoding

$\begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$



Hidden linear layer.  
Effectively a look up for the  
weight row corresponding  
to the input node.

$$\begin{bmatrix} & & & & & \\ & & & & & \\ & & & & & \\ & & & & & \\ & & & & & \\ & & & & & \end{bmatrix} \times \begin{bmatrix} & & & & & \\ & & & & & \\ & & & & & \\ & & & & & \\ & & & & & \\ & & & & & \end{bmatrix} = \begin{bmatrix} & & & & & \\ & & & & & \\ & & & & & \\ & & & & & \\ & & & & & \\ & & & & & \end{bmatrix}$$



Softmax output  
represents probabilities  
of encountering each  
node in a random walk

**Training with Negative Sampling**

- Assume we've sampled many random walks
- From those we can generate sample node pairs  $(u,v)$
- Computing normalized probabilities over all nodes is expensive
- Instead, for each sample, the input is  $u$  and the target output is a 1 hot encoding where the value at the index for node  $v$  is set to 1 and 0 for indices corresponding to a small number of randomly selected nodes

	$T=1$	$T=2$
ABF	AB BF	AF
ABACE	AB BA AC CE	AA BC AE
ABACD	AB BA AC CD	AA BC AD

<sup>†</sup>Grover, Aditya, and Jure Leskovec. "node2vec: Scalable feature learning for networks." Proceedings of the 22nd ACM SIGKDD international conference on Knowledge discovery and data mining. 2016.

# Message passing methods

- Intended to address limitations of shallow embeddings - especially inability to use node feature information and transductive nature (i.e., can only learn embeddings for nodes present during training)
- Methods are motivated by the need to satisfy *permutation invariance*, i.e., the results should not depend on node ordering as there is no natural ordering of nodes in a graph
- Defining characteristic of most recent graph neural networks (GNNs) is the use of *message passing* – vector messages sent between nodes



# Convolutional graph neural networks

- Convolutional in the sense that parameters are shared across the graph
- Assume node features are represented by  $N \times D$  matrix  $X$  ( $N$  nodes,  $D$  features). For graphs where nodes have no features, set  $X = I$
- Every layer,  $H^{(l)}$ , can be written as

$$H^{(l+1)} = f(H^{(l)}, A)$$

where  $A$  is the adjacency matrix,  $H^{(0)} = X$ .

- For *vanilla* convolutional neural network

$$f(H^{(l)}, A) = \sigma(\hat{D}^{-\frac{1}{2}} \hat{A} \hat{D}^{-\frac{1}{2}} H^{(l)} W^{(l)})$$

Diagonal node degree  
matrix (used to  
normalize rows of  $A$ )

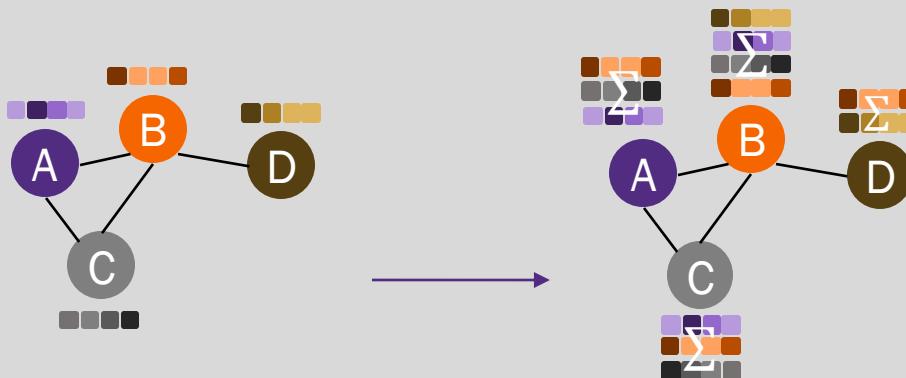
Augmented adjacency  
matrix (diagonal  
elements all set to 1)

Weight matrix for  
layer  $l$ . Importantly,  
the dimension of  $W$  is  
independent of the  
size of the graph.

# Convolutional graph neural networks

After the first layer, each node representation is now the weighted sum of its input feature representation and its neighbors input features (**the messages**)

After the second layer, each node representation is now the weighted sum of its input feature representation and its 1 and 2 hop neighbors

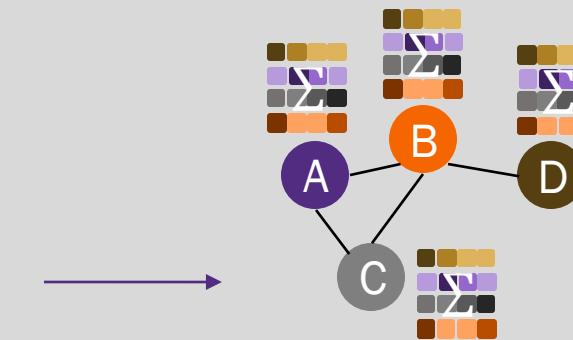


$$H^{(0)} = X \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \\ 10 & 11 & 12 \end{bmatrix}$$

$$H^{(1)} = \sigma(\hat{D}^{-\frac{1}{2}} \hat{A} \hat{D}^{-\frac{1}{2}} X W^{(1)})$$

$$\hat{A} = \begin{array}{c|ccccc} & A & B & C & D \\ \hline A & 1 & 1 & 1 & 0 \\ B & 1 & 1 & 1 & 1 \\ C & 1 & 1 & 1 & 0 \\ D & 0 & 1 & 0 & 1 \end{array}$$

Consider  $\hat{A}X$ . This updates the feature row for each node as the sum of its features and its neighbor's features



$$H^{(2)} = \sigma(\hat{D}^{-\frac{1}{2}} \hat{A} \hat{D}^{-\frac{1}{2}} H^{(1)} W^{(2)})$$

Consider  $\hat{A}H^1$ . This updates the feature row for each node as in the first layer, but now the representation of each node neighbor contains information from their neighbors as well.

# Graph attention

- Similar attention mechanism to transformers. Recall the QKV paradigm.
- The key (K) and value (V) are the transformed message from a neighboring node (e.g., output of an MLP applied to message).
- The receiving node uses the transformed version of its feature representation as the query (Q)



$$\alpha_{ij} = \frac{\exp(\sigma(\mathbf{a}[\mathbf{W}\mathbf{h}_i] || \mathbf{W}\mathbf{h}_j]))}{\sum_{k \in N_i} \exp(\sigma(\mathbf{a}[\mathbf{W}\mathbf{h}_i] || \mathbf{W}\mathbf{h}_k)))}$$

$$h'_i = \sigma \left( \sum_{j \in N_i} \alpha_{ij} \mathbf{W}\mathbf{h}_j \right)$$

Message transformation  
MLP weights

GCN weights

Attention weight applied  
to message from node  $j$   
to node  $i$ , where  $N_i$  are  
the neighbors of node  $i$   
(which includes node  $i$ )

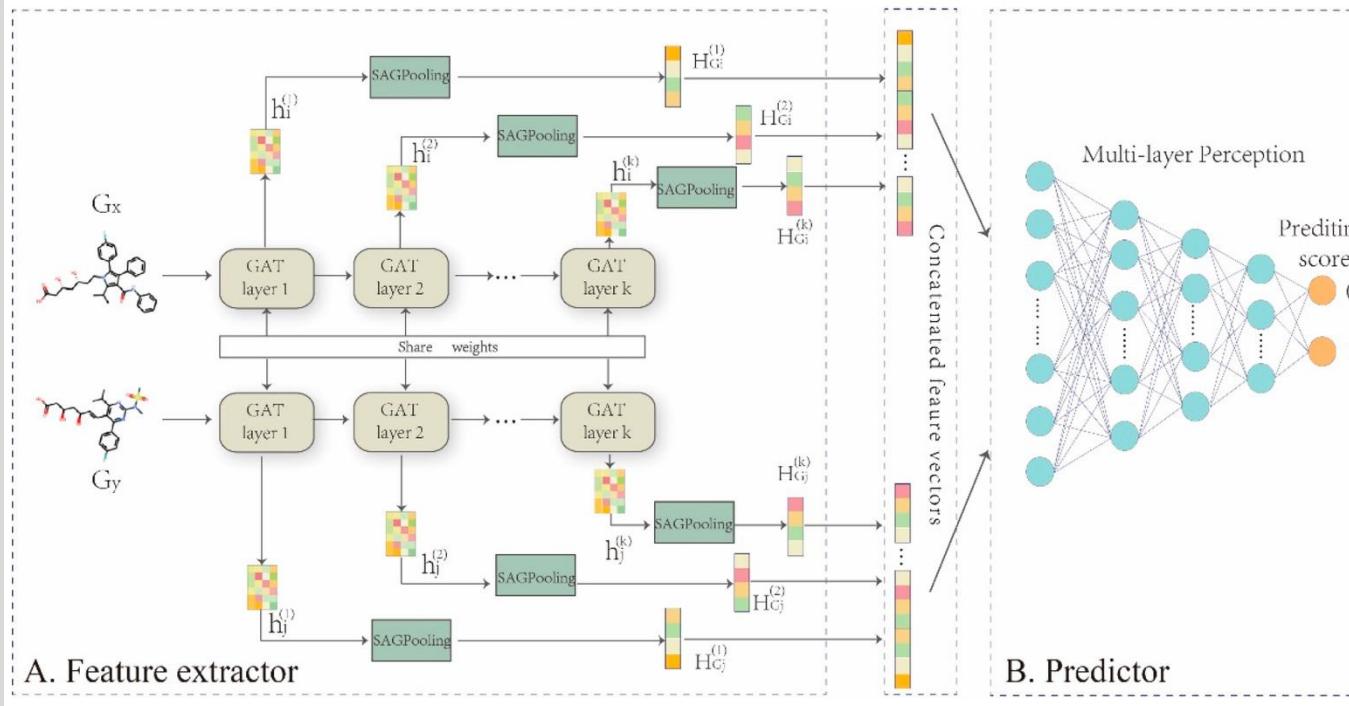
Attention based  
representation of node  $i$



# Applications

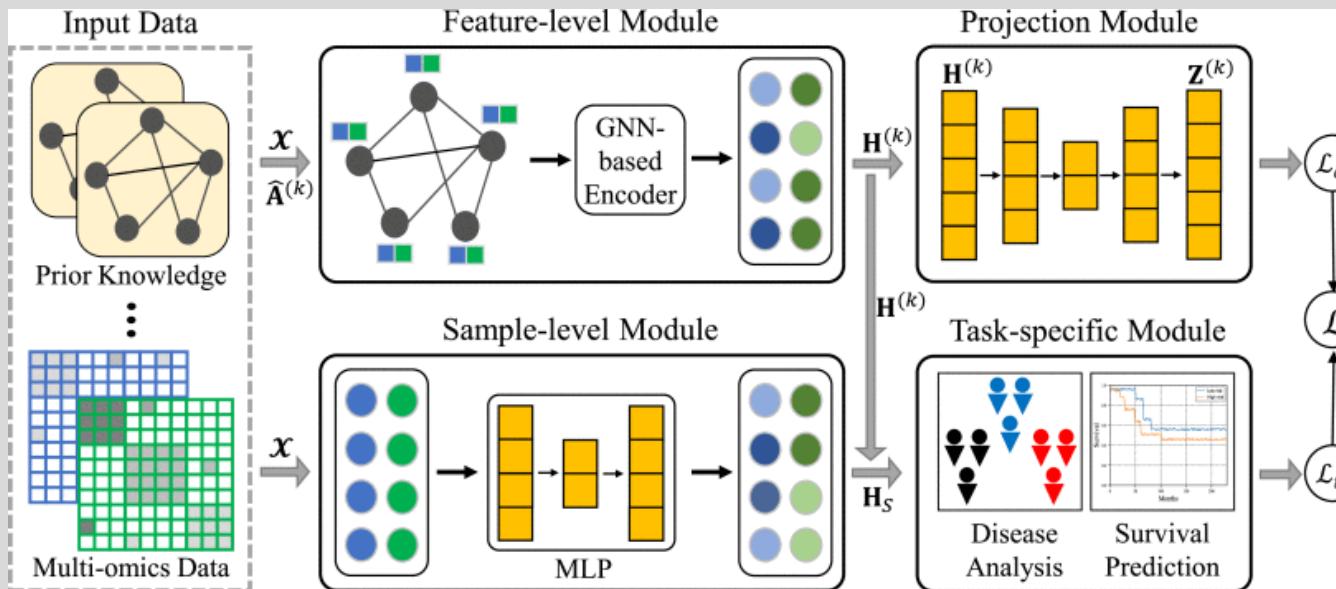


# GNN drug-drug interaction prediction



- Each drug is represented by a graph
  - Nodes represent atoms
  - Edges represent bonds
  - Node features (78 dims) include atomic symbol, valence, formal charge, number of free radical electrons, hybridization, number of hydrogen bonds, and aromaticity
- Each drug in a candidate DDI pair is passed through a 5-layer GCN with attention
  - Output node representations are down sampled via Self-attention graph (SAG) pooling to form an overall representation of each drug
  - Drug representations are concatenated and passed through a FFN for DDI classification

# Integrating prior knowledge for cancer subtype prediction



- Prior knowledge is represented as a graph
  - Nodes represent genes (selected based on omics)
  - Edges represent relationships between genes (bespoke rules based on the knowledge base)
  - Graph representation from GNN encoder
  - Projection model attempts to maximally contrast gene representations from knowledge base
- Patient level omics data processed in MLP
- Output representations of prior knowledge and patient omic data combined and passed to task-specific module

# Biology informed AI – improving explainability

*Central hypothesis:* AI that utilizes abductive reasoning through integration of formal biology knowledge with observational data will demonstrate **improved** inference performance and **explainability in limited data settings.**

## Our Methodological Research

- Selective attention mechanisms to identify knowledge graph concepts relevant to AI system output
- Pathway analysis methods to trace the paths in the knowledge graph that lead to a particular prediction

